

# CS321: Computer Architecture

Course web:

<http://172.16.1.6/~jimson/>

# Computer Architecture

Basic functional blocks of a computer: CPU, memory, input-output subsystems, control unit. Instruction set architecture of a CPU - registers, instruction execution cycle, RTL interpretation of instructions, addressing modes, instruction set. Case study - instruction sets of some common CPUs; Assembly language programming for some processor; Data representation: signed number representation, fixed and floating point representations, character representation. Computer arithmetic - integer addition and subtraction, ripple carry adder, carry look-ahead adder, etc. multiplication – shift-and-add, Booth multiplier, carry save multiplier, etc. Division - non-restoring and restoring techniques, floating point arithmetic; CPU control unit design: hardwired and microprogrammed design approaches, Case study - design of a simple hypothetical CPU; Pipelining: Basic concepts of pipelining, throughput and speedup, pipeline hazards; Memory organization: Memory interleaving, concept of hierarchical memory organization, cache memory, cache size vs block size, mapping functions, replacement algorithms, write policy; Peripheral devices and their characteristics: Input-output subsystems, I/O transfers - program controlled, interrupt driven and DMA, privileged and non-privileged instructions, software interrupts and exceptions. Programs and processes - role of interrupts in process state transitions.

# Prerequisite & Text Book

- Prerequisite: **CS225 & CS226**
- Text Books
  - David A. Patterson, John L. Hennessy, Computer Organization and Design, Fourth Edition:
  - The Hardware/Software Interface, Morgan Kaufmann; 4 edition, 2011.
  - 2. A. Tenenbaum, Structured Computer Organization, 4th Ed, Prentice-Hall of India, 1999.
  - 3. W. Stallings, Computer Organization and Architecture: Designing for Performance, 6th Ed, Prentice Hall, 2005.
  - 4. J. Hennessy and D. Patterson, Computer Architecture A Quantitative Approach, 3rd Ed, Morgan Kaufmann, 2002..

# Evaluation Policy(CS321)

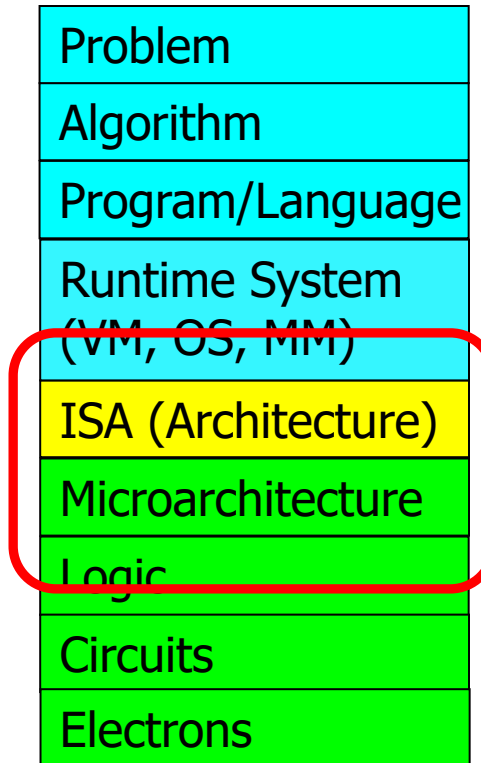
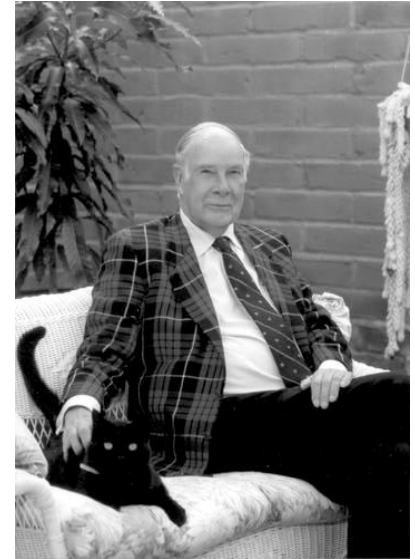
- Assignments  
& Quizzes : 40%
- Mid-Sem : 25%
- End-Sem : 25%
- Performance:10 %

# Architecture Lab(CS322)

- Regular Lab work, 40%
- Mid-Sem : 15%
- End-Sem : 25%
- Mini Project:45 %

# Levels of Transformation

“The purpose of computing is insight” (*Richard Hamming*)  
*We gain and generate insight by solving problems*  
*How do we ensure problems are solved by electrons?*



# The Power of Abstraction

- Levels of transformation create abstractions
  - Abstraction: A higher level only needs to know about the interface to the lower level, not how the lower level is implemented
  - E.g., high-level language programmer does not really need to know what the ISA is and how a computer executes instructions
- Abstraction improves productivity
  - No need to worry about decisions made in underlying levels
  - E.g., programming in Java vs. C vs. assembly vs. binary vs. by specifying control signals of each transistor every cycle
- Then, why would you want to know what goes on underneath or above?

# Why CS321?

- What if

- The program you wrote is running slow?
- The program you wrote does not run correctly?
- The program you wrote consumes too much energy?

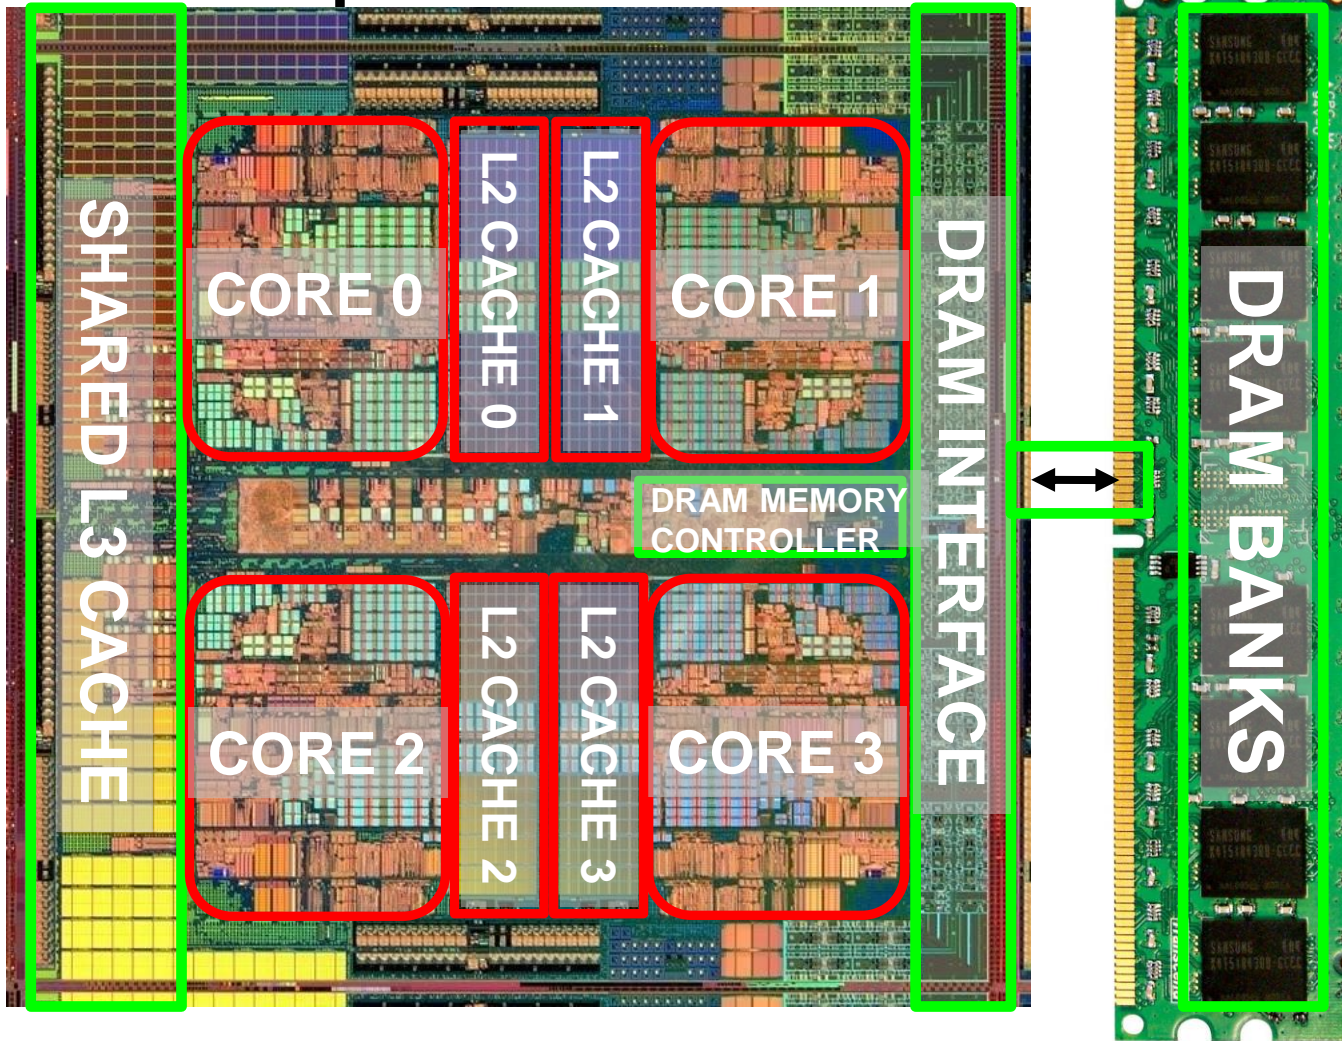
- What if

- The hardware you designed is too hard to program?
- The hardware you designed is too slow because it does not provide the right primitives to the software?



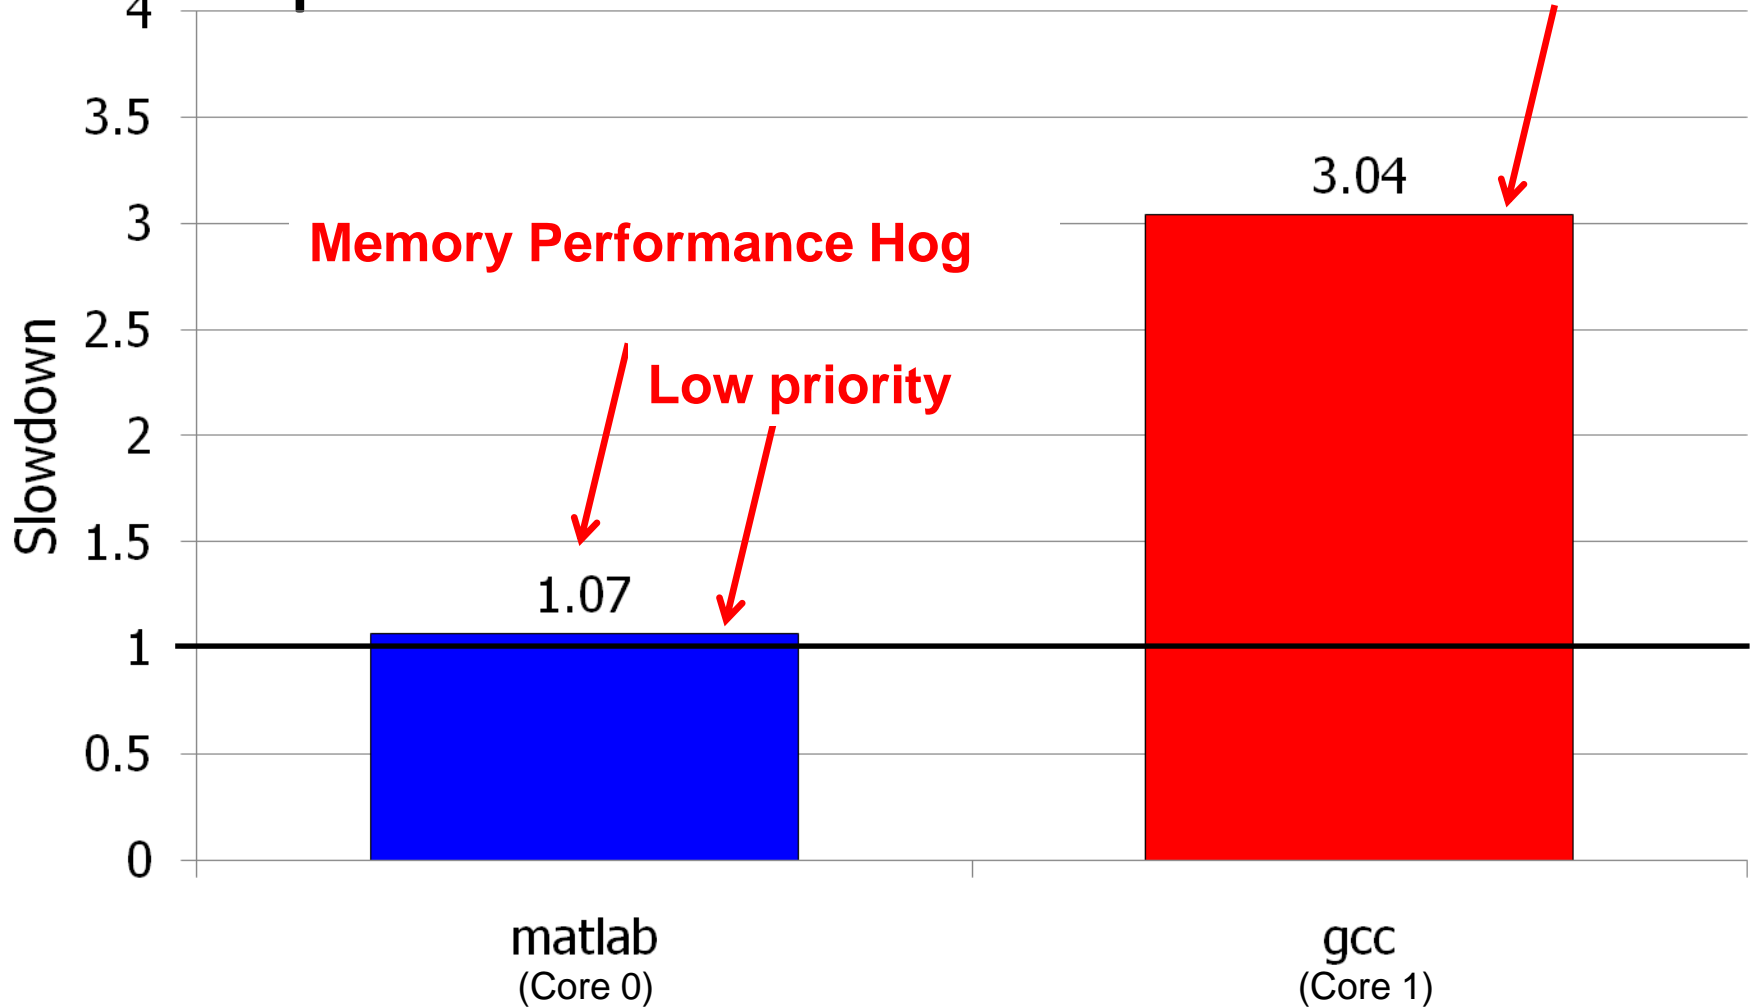
# An Example: Multi-Core Systems

Multi-Core  
Chip



\*Die photo credit: AMD Barcelona

# Unexpected Slowdowns in Mul

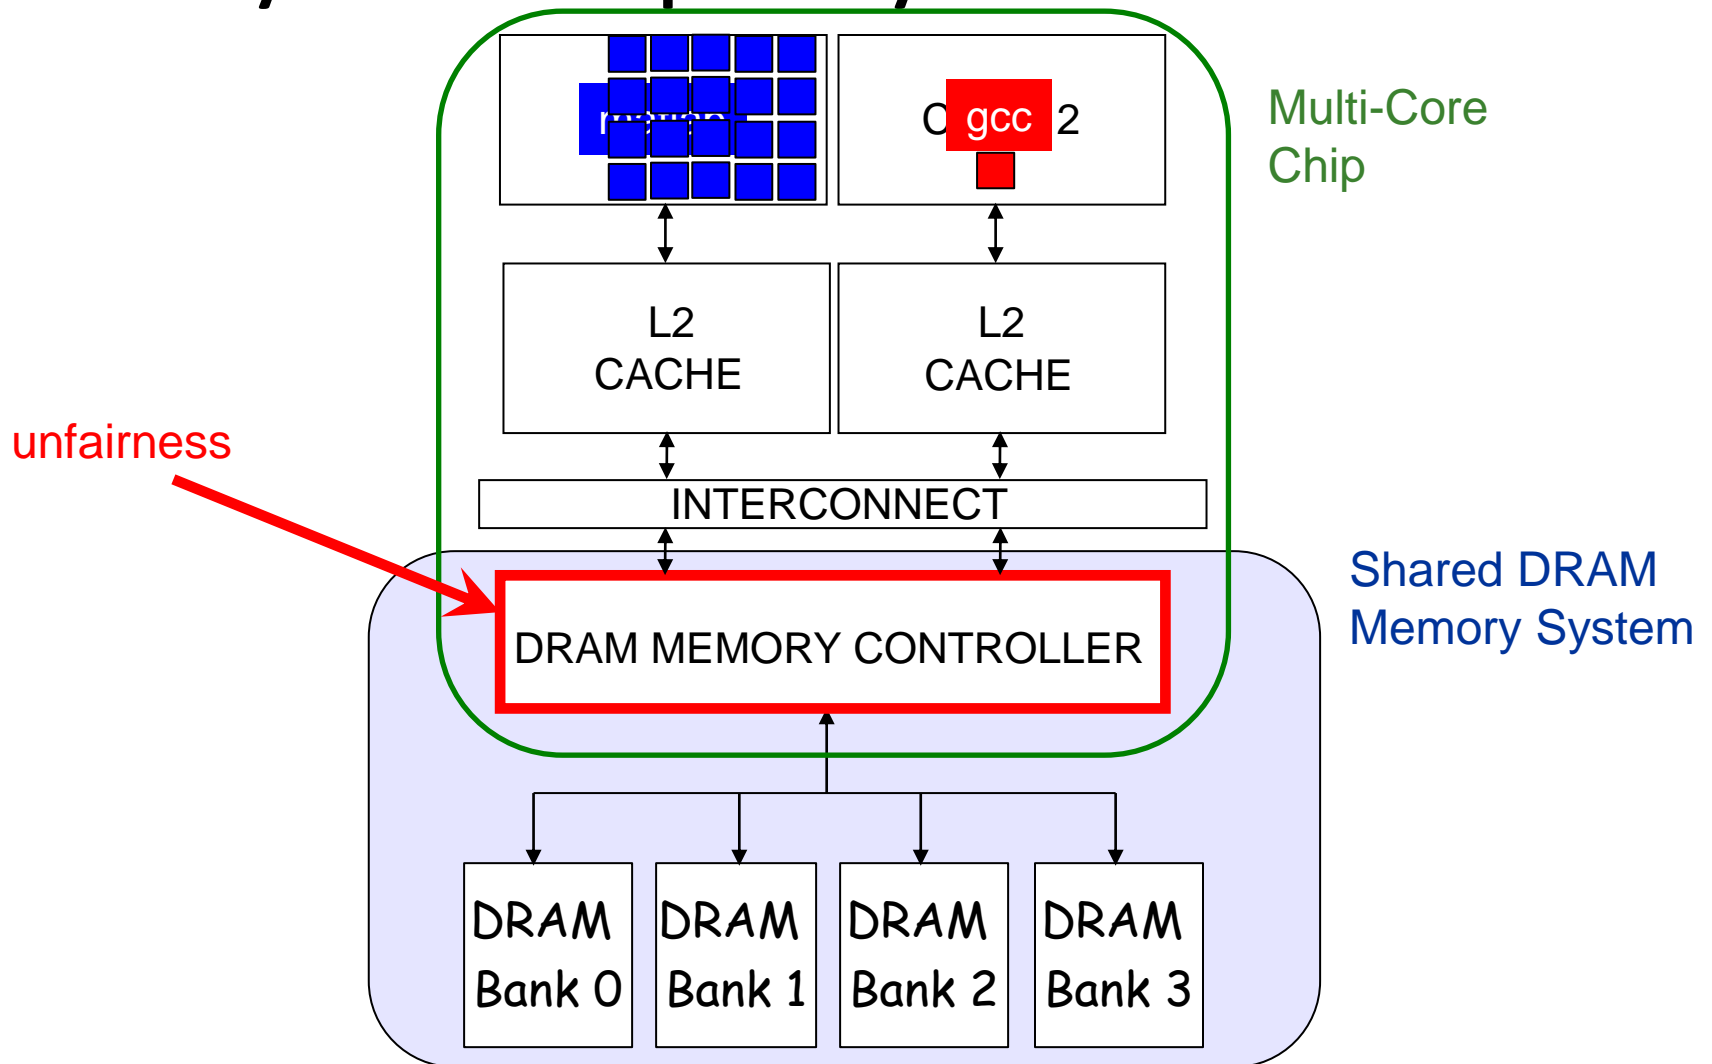


Moscibroda and Mutlu, “[Memory performance attacks: Denial of memory service in multi-core systems](#),” USENIX Security 2007.

## A Question or Two

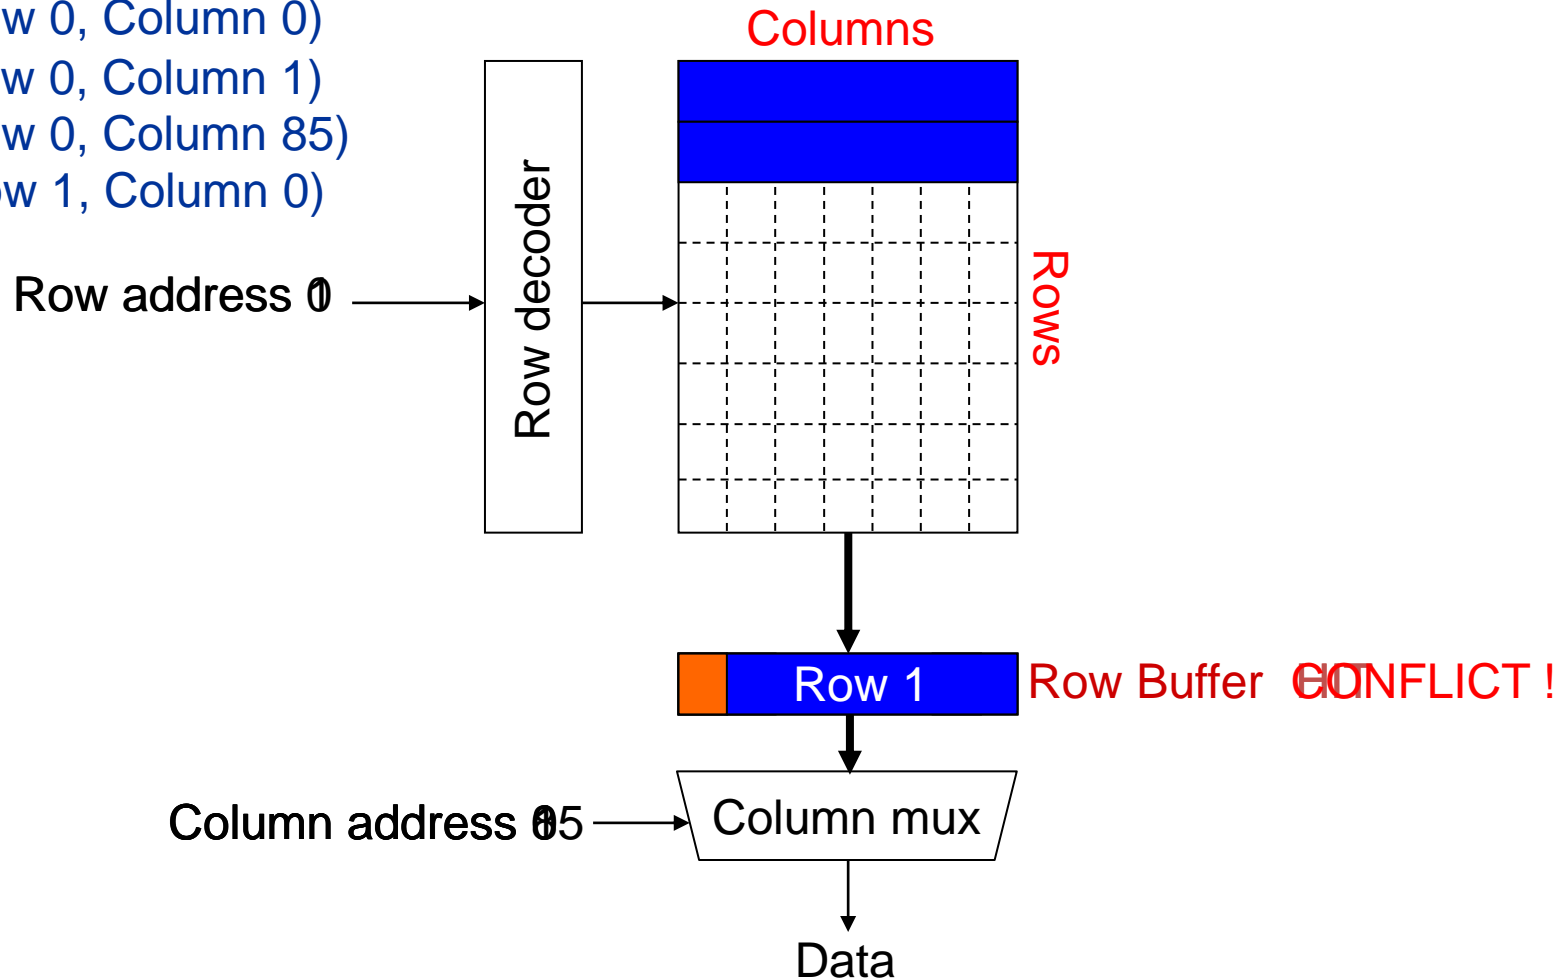
- Can you figure out why there is a disparity in slowdowns if you do not know how the processor executes the programs?
- Can you fix the problem without knowing what is happening “underneath”?

# Why the Disparity in Slowdowns?



# DRAM Bank Operation

Access Address:  
(Row 0, Column 0)  
(Row 0, Column 1)  
(Row 0, Column 85)  
(Row 1, Column 0)



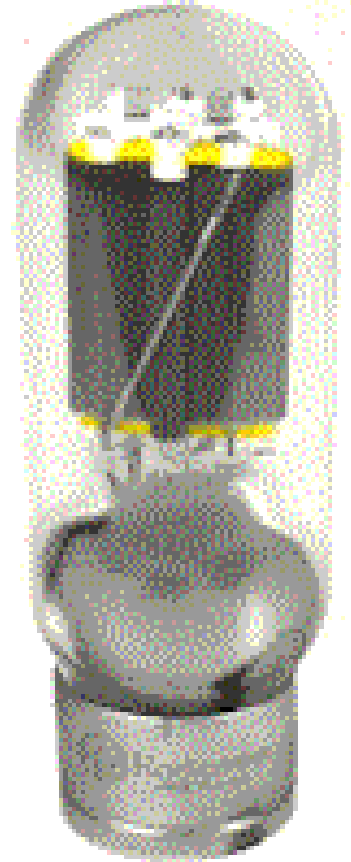
# INTRODUCTION

# Generations of Computer

1. **First Generation - 1940-1956: Vacuum Tubes**
2. **Second Generation - 1956-1963: Transistors**
3. **Third Generation - 1964-1971: Integrated Circuits**
4. **Fourth Generation - 1971-Present: Microprocessors**
5. **Fifth Generation - Present and Beyond: Artificial Intelligence**

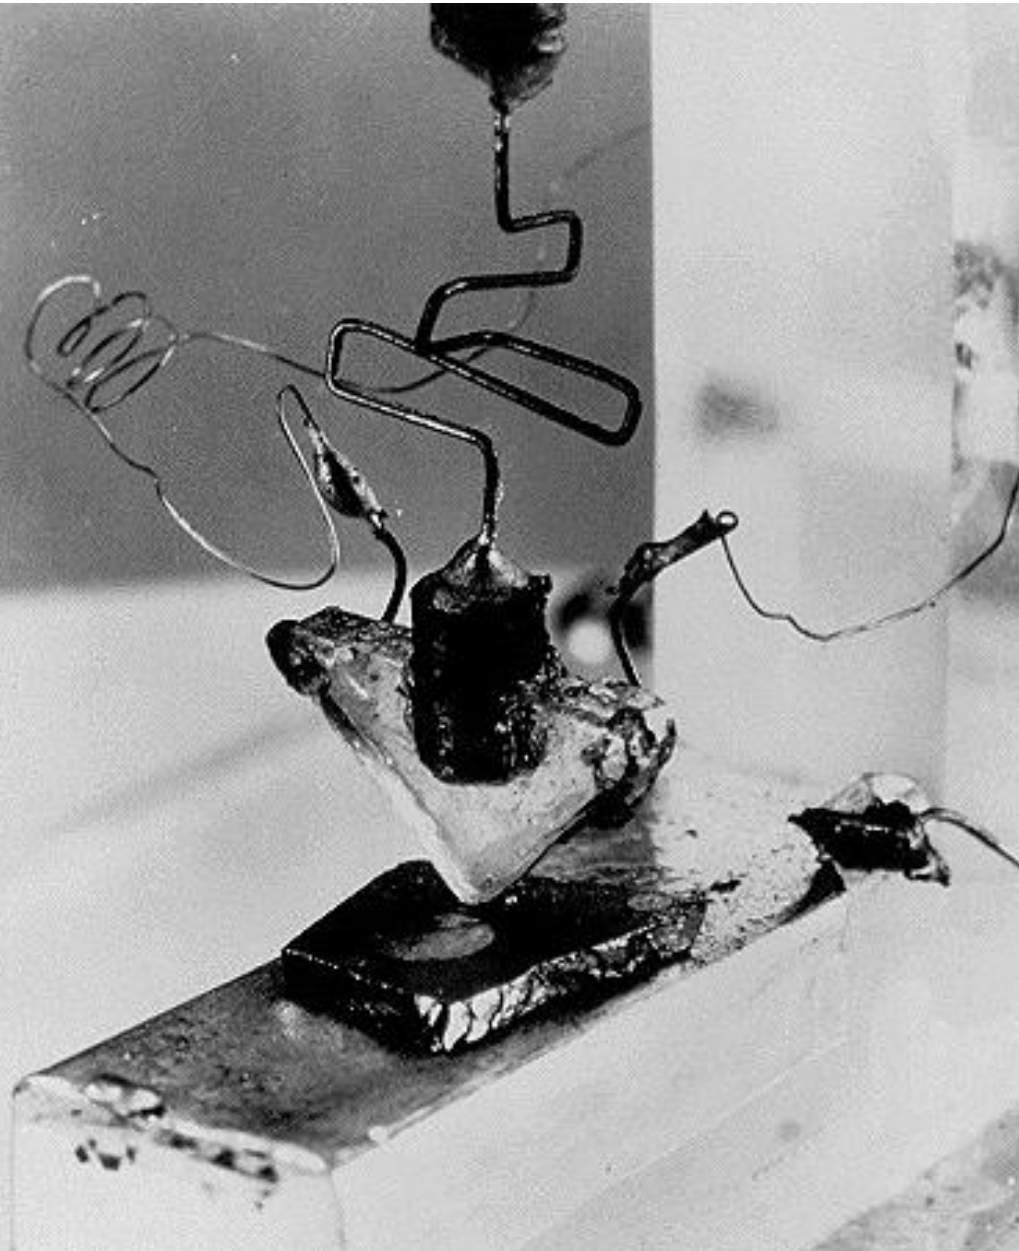
# FIRST GENERATION 1940 - 1956

- First generation computers used Vacuum Tubes
- Vacuum tubes are glass tubes with circuits inside.
- The word vacuum indicates that they have no air inside, which protects the circuitry.
- Building a computer with these vacuum tubes would result in a very large machine occupying one full room.



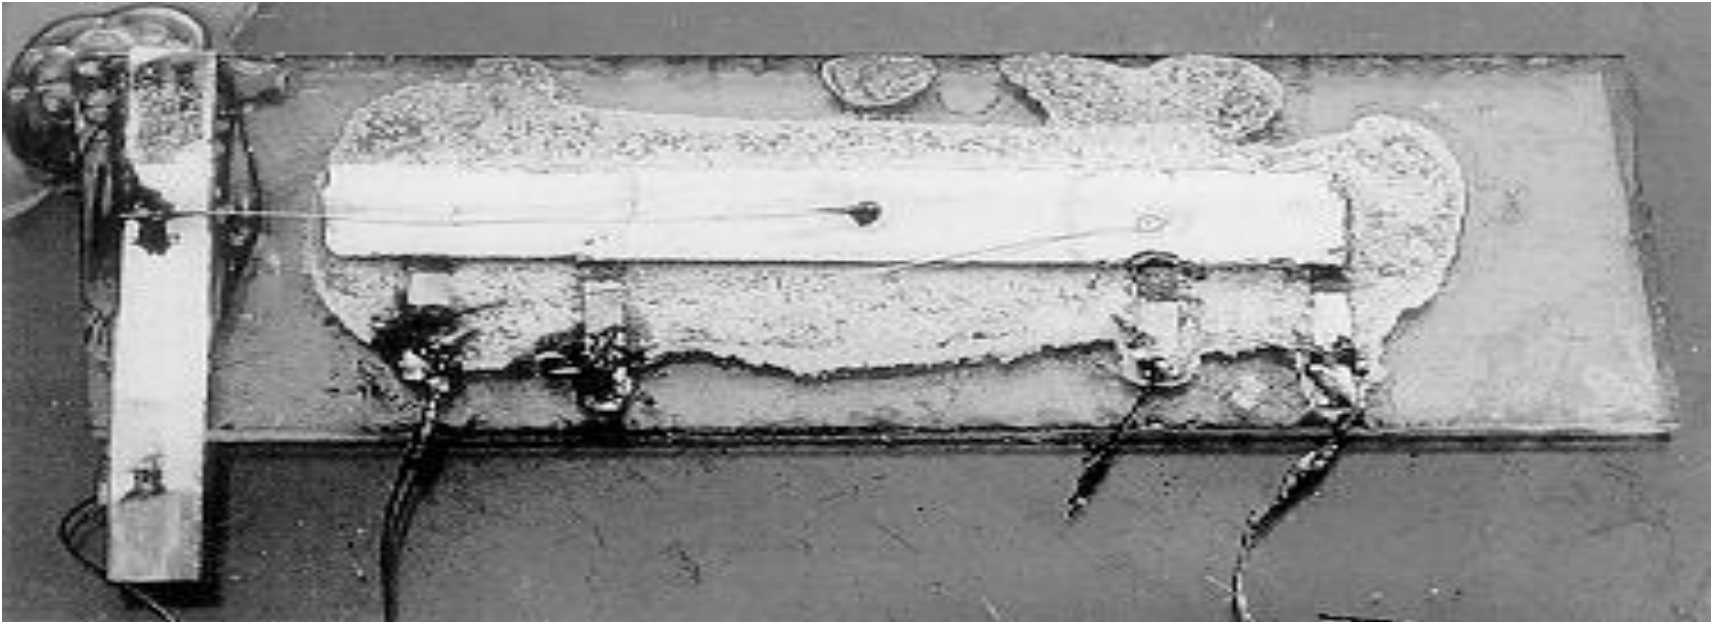


# First Transistor



- Uses Silicon
- developed in 1948
- won a Nobel prize
- on-off switch
- Second Generation Computers used Transistors, starting in 1956

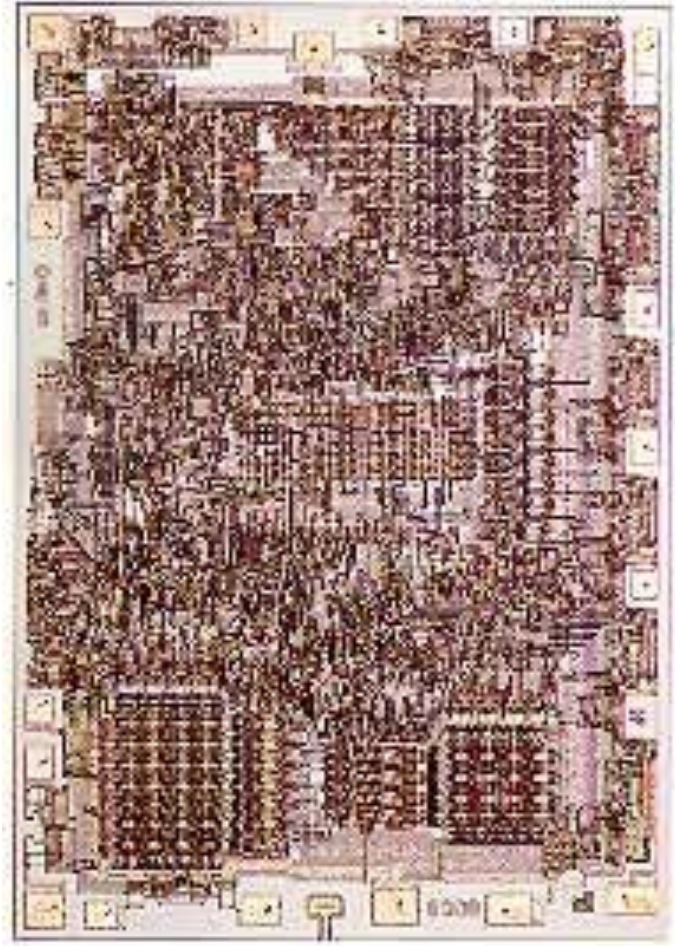
# Integrated Circuits



- Third Generation Computers used Integrated Circuits (chips).
- Integrated Circuits are transistors, resistors, and capacitors integrated together into a single "chip"
- Instead of punched cards and printouts, users started interacting with keyboards and mouse.

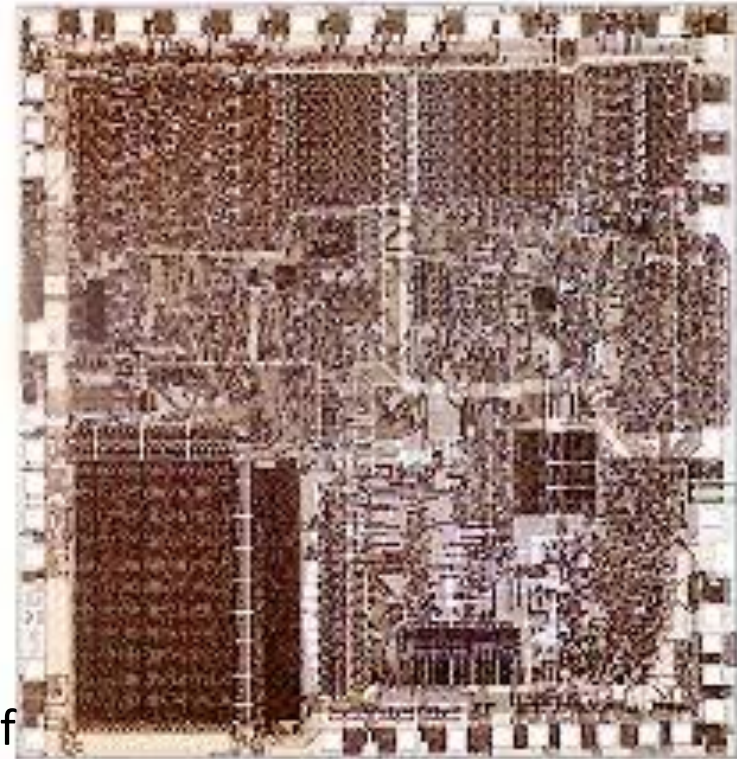
## 1972: 8008 Microprocessor

- The 8008 was twice as powerful as the 4004.
- According to the magazine *Radio Electronics*, Don Lancaster, a dedicated computer hobbyist, used the 8008 to create a predecessor to the first personal computer, a device *Radio Electronics* dubbed a "TV typewriter." It was used as a dumb terminal.



# 1978: 8086-8088 Microprocessor

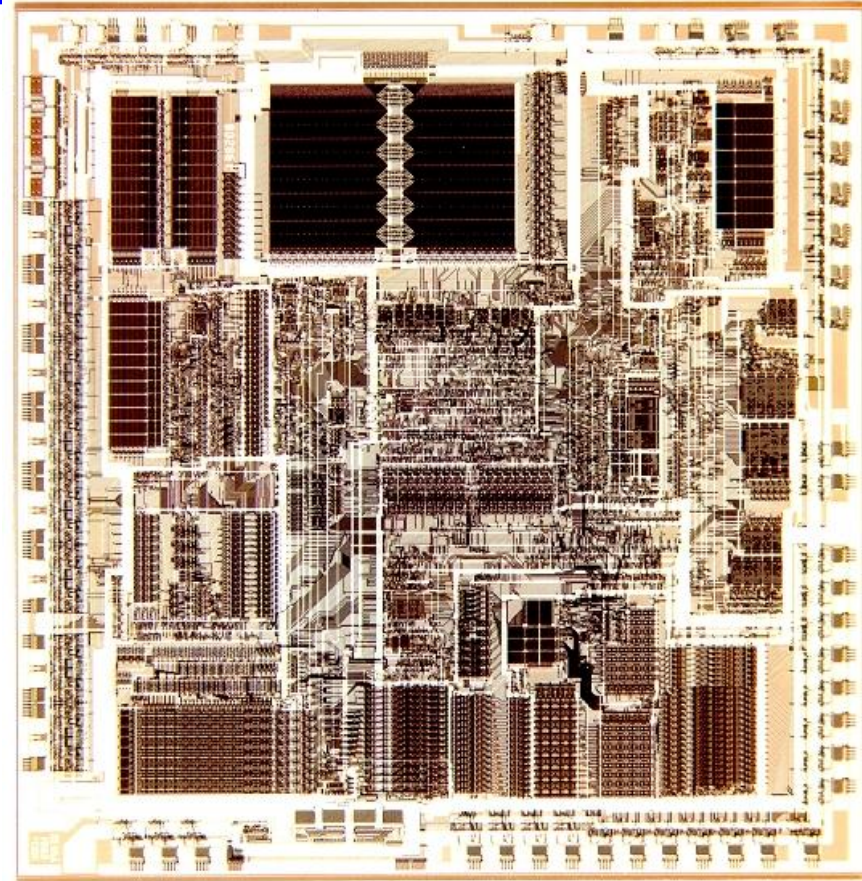
- In 1978 intel
  - 16-bit
  - 2.5MIPS **millions of instructions per second**
  - **1M-byte memory**
  - 4- or 6-byte instruction (cache) queue that prefetch instructions
  - over 20,000 variations instructions.
  - 64K bytes of memory found in 8-bit microprocessors to execute efficiently
  - The 16-bit 8086 and 8088 provided 1M byte of memory for these applications
  - **Popularity of Intel ensured in 1981 when IBM chose the 8088 in its personal computer**





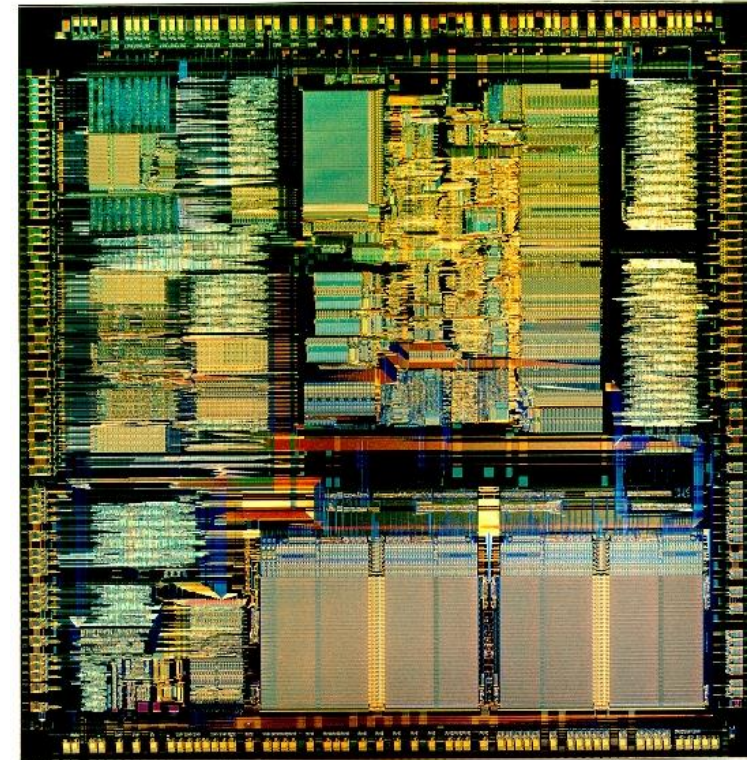
# 1982: 286 Microprocessor

- The 286, also known as the 80286, was the first Intel processor that could run all the software written for its predecessor.
- This software compatibility remains a hallmark of Intel's family of microprocessors.
- Within 6 years of its release, there were an estimated 15 million 286-based personal computers installed around the world.
- 80286: updated 8086
  - 16M byte memory addressing
  - Instructions identical to 8086 few more added
  - 4MIPS
  - 8Mhz clock speed



# 1985: Intel 386(TM) Microprocessor

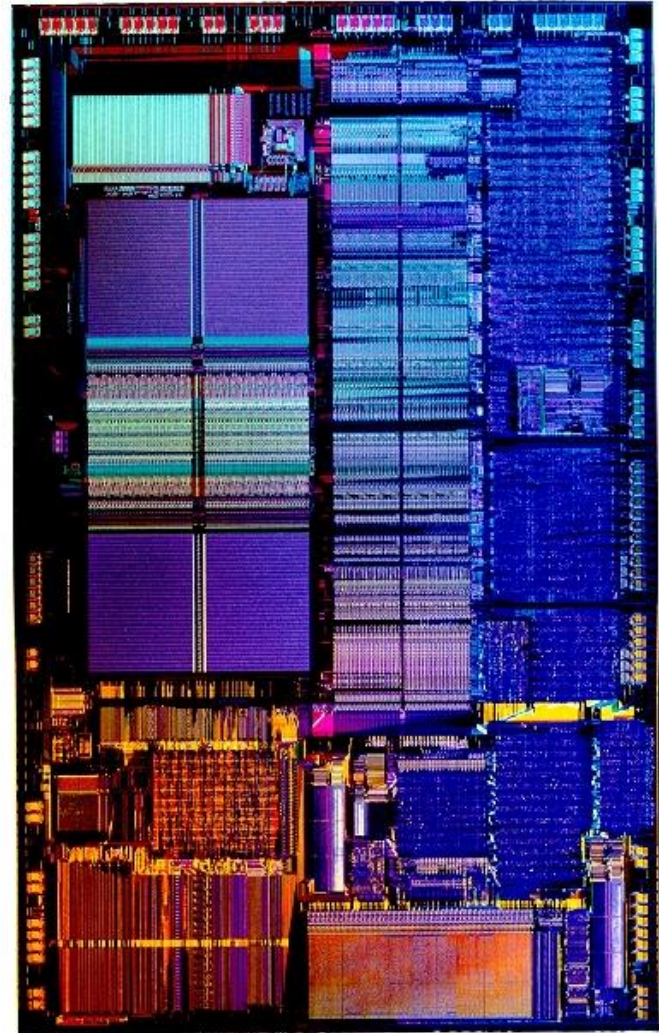
- The Intel 386™ microprocessor featured 275,000 transistors--more than 100times as many as the original 4004.
- It was a 32-bit chip and was "multi tasking," meaning it could run multiple programs at the same time.
- 80386 (1985)
  - 32-bit address bus and 32-bit data bus
  - 4GB memory
  - Hardware circuitry for memory management
  - Additional instructions referenced 32-bit registers and managed the memory system





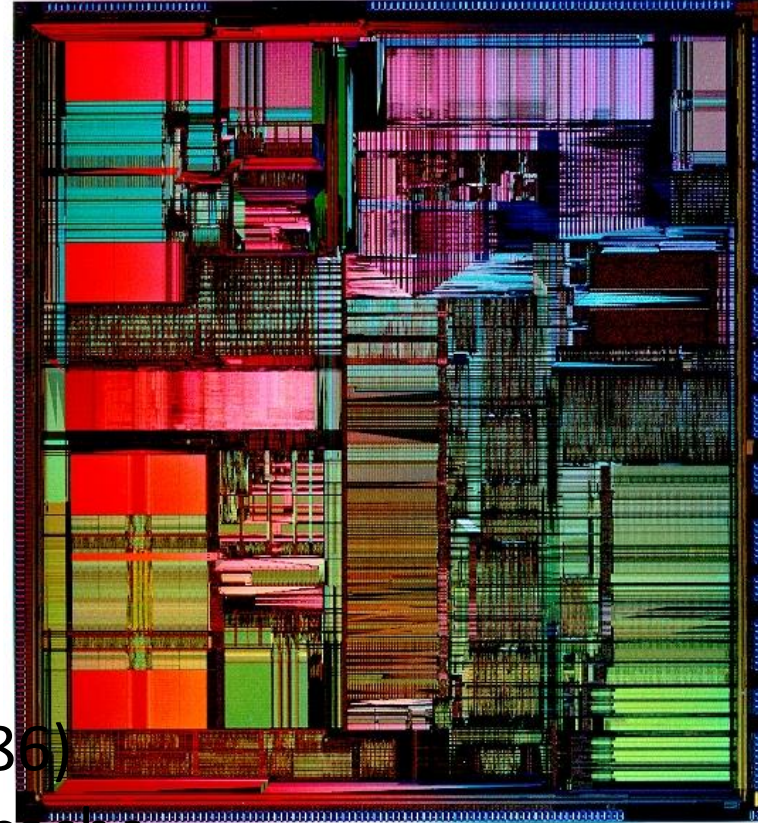
# 1989: Intel 486(TM) DX CPU

- The 486<sup>TM</sup> generation really allowed the shift from command-level computing into point-and-click computing.
- The Intel 486<sup>TM</sup> processor was the first to offer a built-in math coprocessor, which speeds up computing because it offers complex math functions from the central processor.
- 80486
  - Highly integrated package.
  - 80386-like microprocessor.80387-like numeric coprocessor.
  - 50 MIPS
  - 8K-byte cache memory system
  - Half of its instructions executed in 1 clock cycle rather 2



# 1993: Pentium® Processor

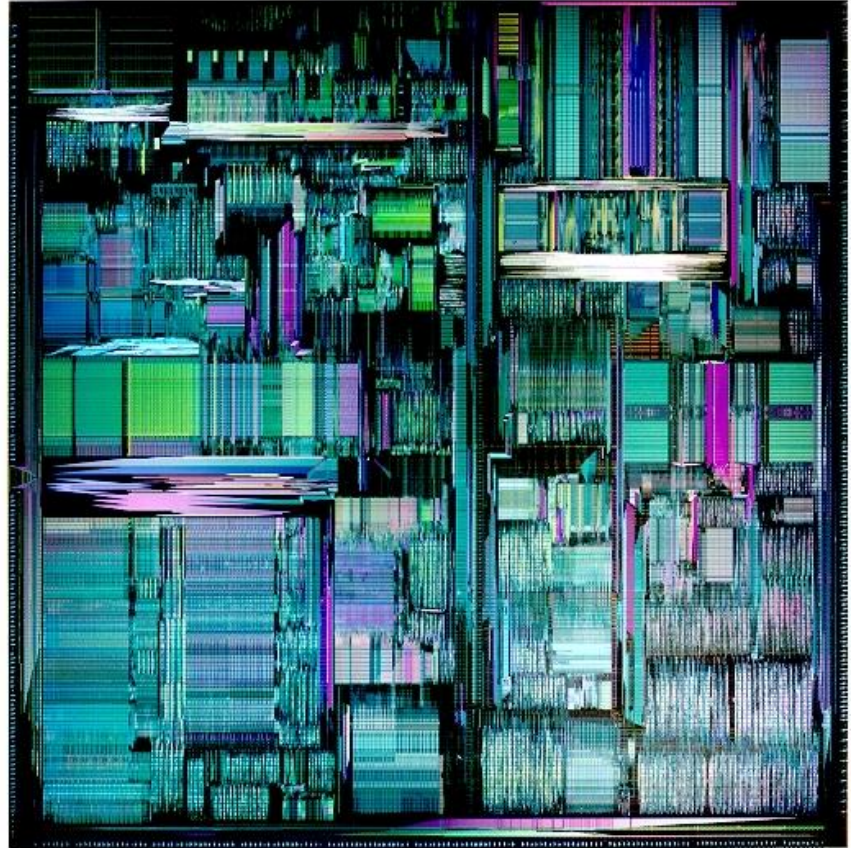
- The Pentium® processor allowed computers to more easily incorporate "real world" data such as speech, sound, handwriting and photographic images.
- The name Pentium®, mentioned in the comics and on television talk shows, became a household word soon after introduction.
  - Originally named P5 or 80586
  - Clock speed of 60MHz
  - Executes 110MIPS
- Cache size: 16K bytes (8K cache in 80486)
  - 8K-byte instruction cache and data cache.
- Memory system up to 4G bytes.
- Data bus width increased to a full 64 bits.
- Data bus transfer speed 60 MHz or 66 MHz.
  - depending on the version of the Pentium





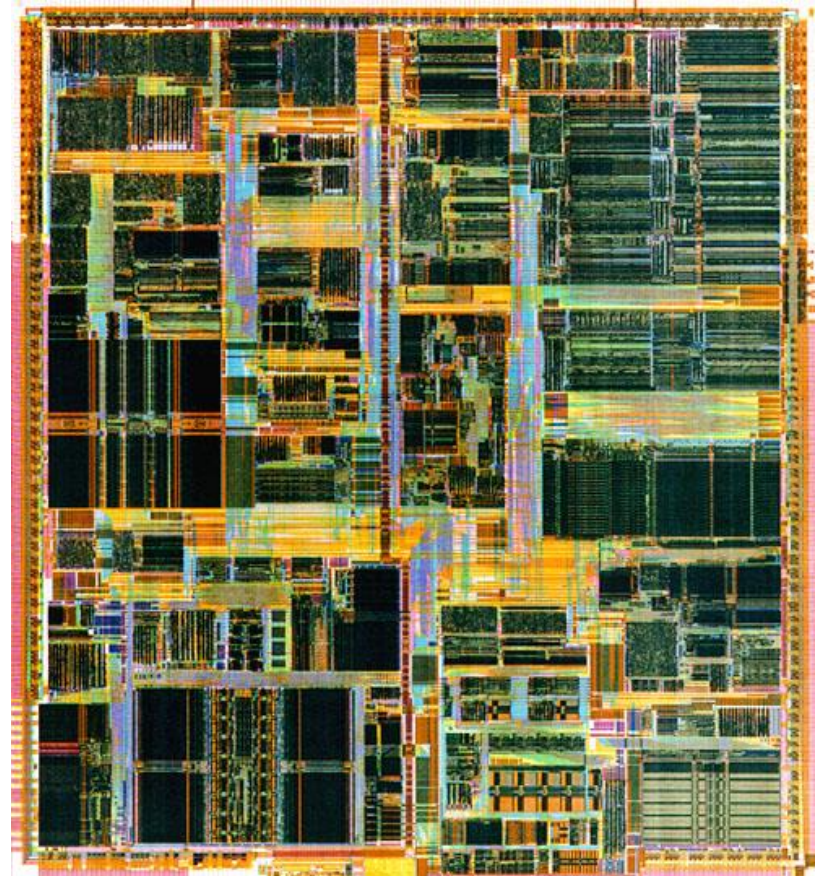
# 1995: Pentium® Pro Processor

- Released in the fall of 1995 the Pentium® Pro processor is designed to fuel 32-bit server and workstation-level applications, enabling fast computer-aided design, mechanical engineering and scientific computation.
- Each Pentium® Pro processor is packaged together with a second speed-enhancing cache memory chip.
- The powerful Pentium® Pro processor boasts 5.5 million transistors.



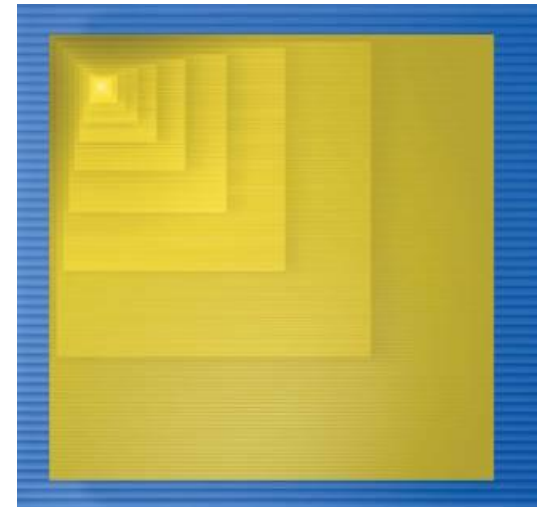
## 1997: Pentium® II Processor

- The 7.5 million-transistor Pentium® II processor incorporates Intel MMX™ technology, which is designed specifically to process video, audio and graphics data efficiently.
- It is packaged along with a high-speed cache memory chip in an innovative Single Edge Contact (S.E.C.) cartridge that connects to a motherboard via a single edge connector, as opposed to multiple pins.
- With this chip, PC users can capture, edit and share digital photos with friends and family via the Internet; edit and add text, music or between-scene transitions to home movies; and, with a video phone, send video over standard phone lines and the Internet.



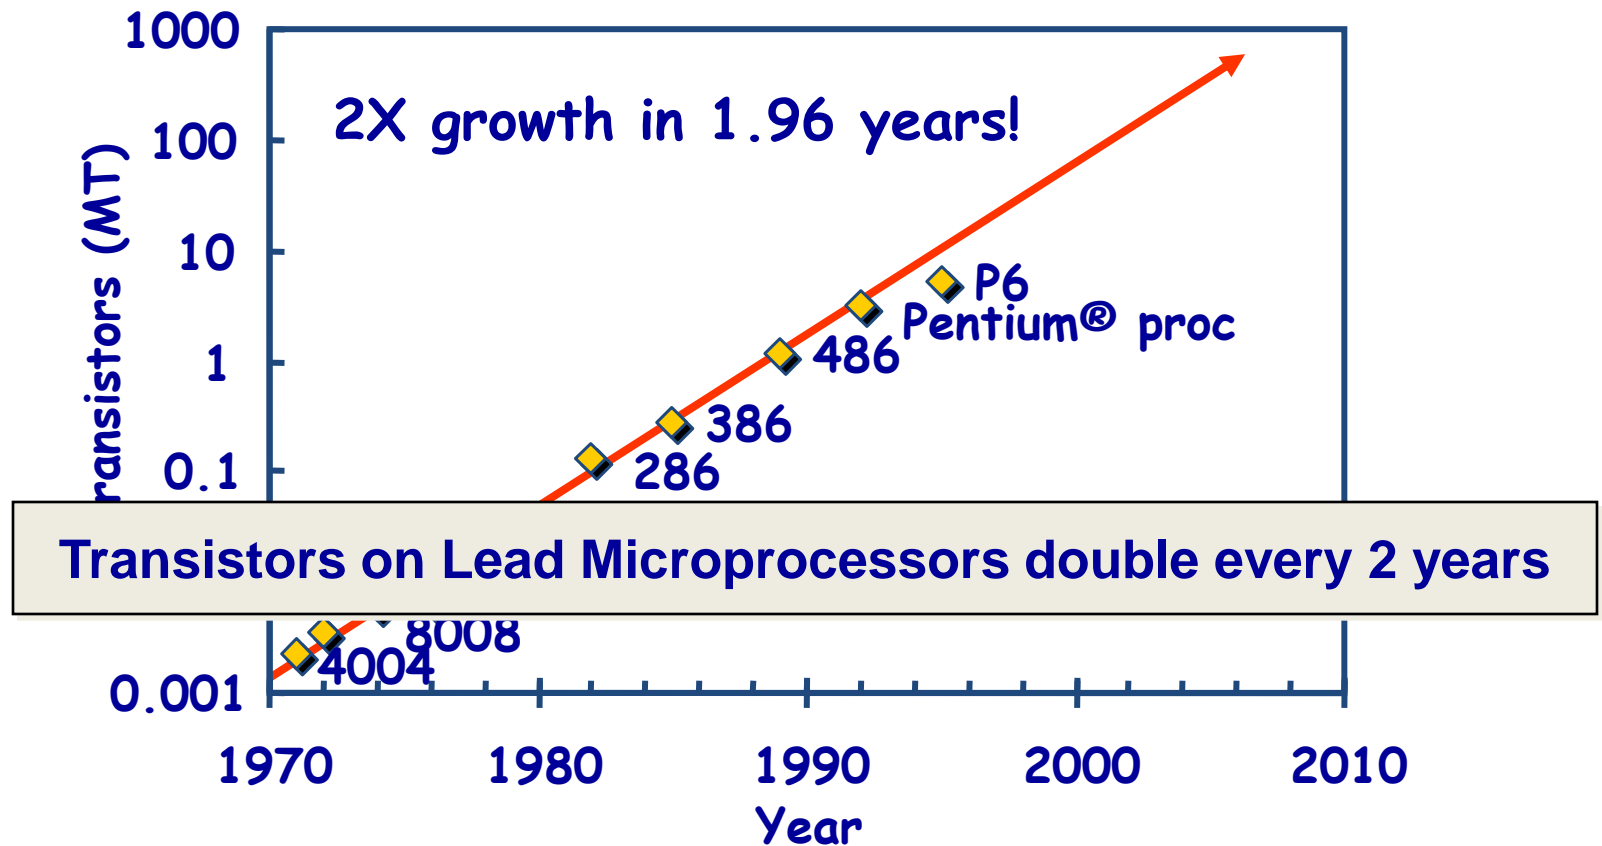
# Moore's Law

- IC capacity doubling about every 18 months for several decades
  - Known as “Moore’s Law” after Gordon Moore, co-founder of Intel
    - Predicted in 1965 predicted that components per IC would double roughly every year or so
  - Picture depicts related phenomena
    - For a particular number of transistors, the IC shrinks by half every 18 months
      - Notice how much shrinking occurs in just about 10 years
      - Enables incredibly powerful computation in incredibly tiny devices
  - Today’s ICs hold *billions* of transistors
    - The first Pentium processor (early 1990s) needed only 3 million



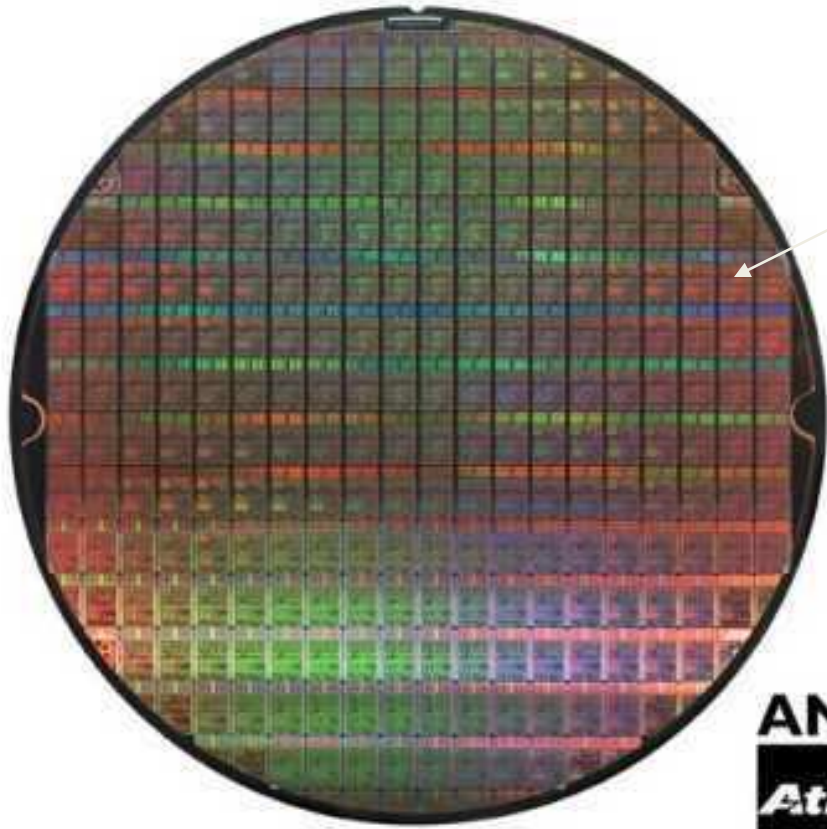


# Moore's law in Microprocessors



Courtesy, Intel

# Die



Single die

Wafer



Going up to 12" (30cm)

From <http://www.amd.com>

# *In Picture...*

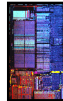


**Silicon Process Technology** 1.5 $\mu$  1.0 $\mu$  0.8 $\mu$  0.6 $\mu$  0.35 $\mu$  0.25 $\mu$  0.18 $\mu$  0.13 $\mu$

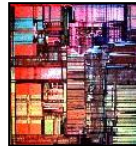
**Intel386™ DX  
Processor**



**Intel486™ DX  
Processor**



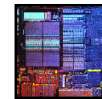
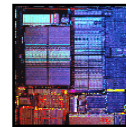
**Pentium®  
Processor**



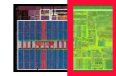
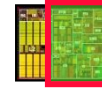
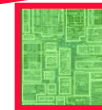
**Pentium® Pro  
Processor**



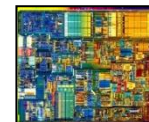
**Pentium® II  
Processor**



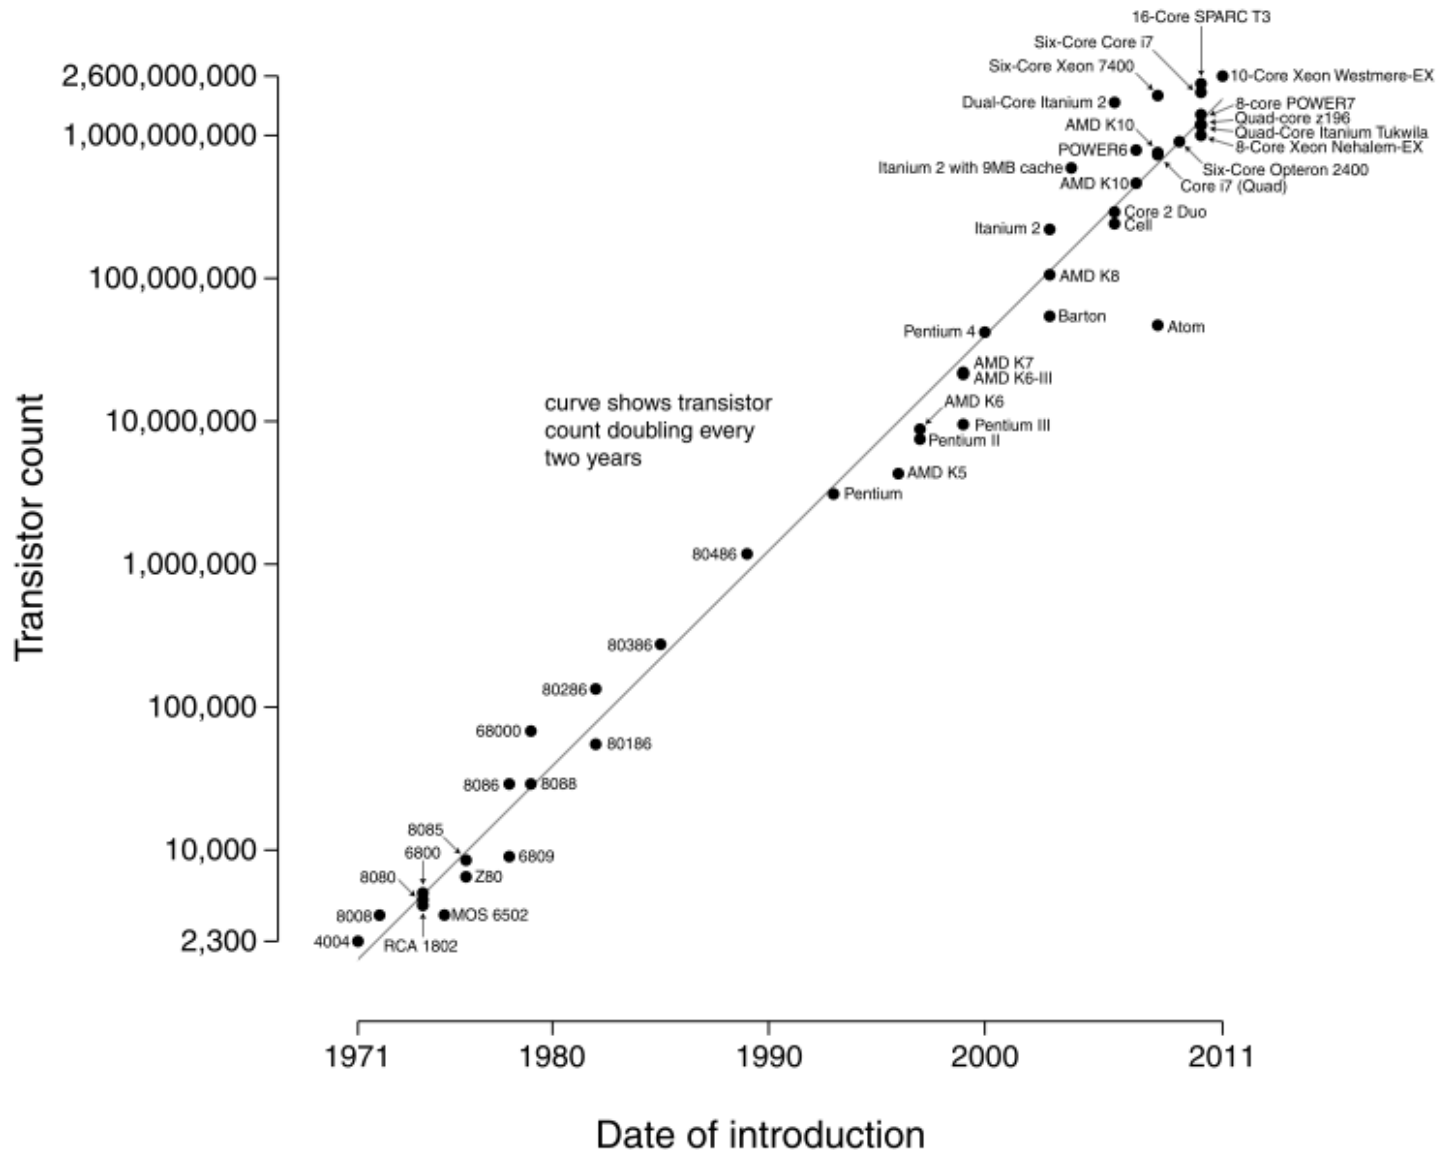
**Pentium® III  
Processor**



**Pentium® 4  
Processor**

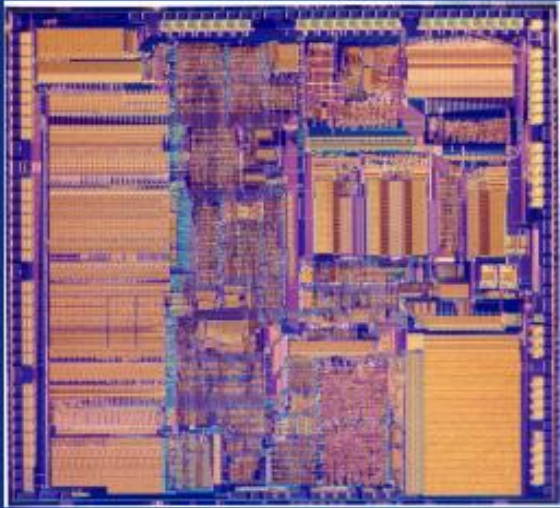


# Microprocessor Transistor Counts 1971-2011 & Moore's Law



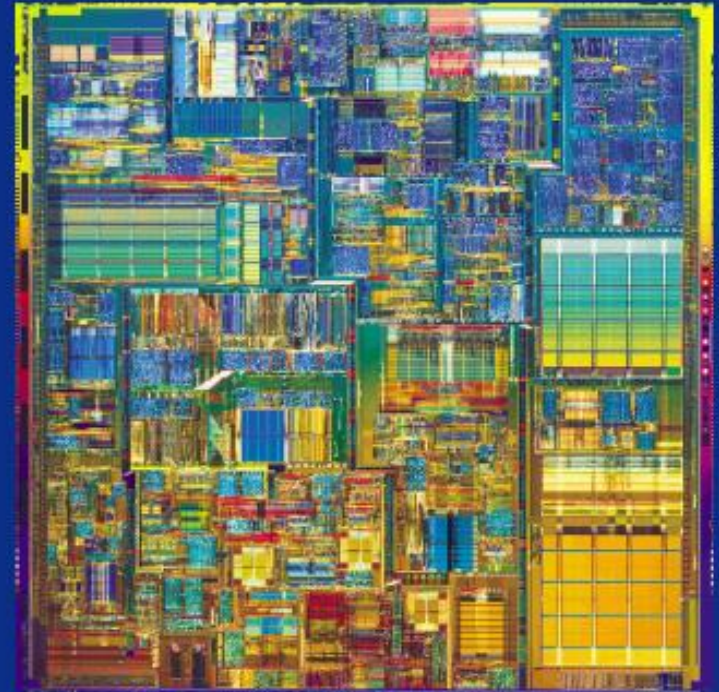
Number of transistors on an integrated circuit doubles ~ every two years

# Scaling at its best



**386 Processor**

May 1986  
@16 MHz core  
275,000 1.5 $\mu$  transistors  
~1.2 SPECint2000



**Pentium® 4 Processor**

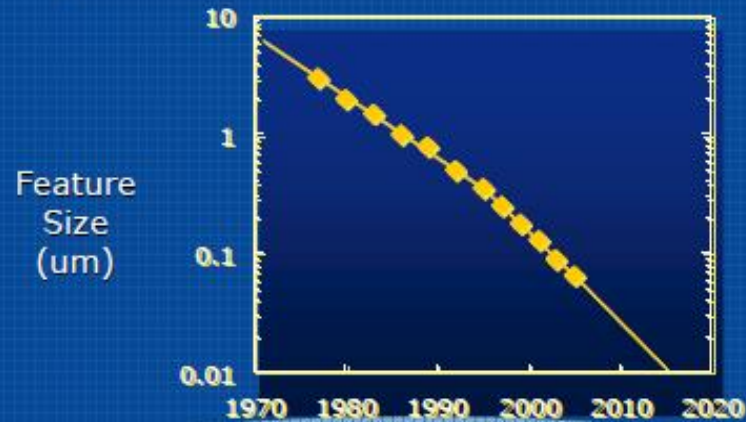
17 Years  
200x  
200x/11x  
1000x

August 27, 2003  
@3.2 GHz core  
55 Million 0.13 $\mu$  transistors  
1249 SPECint2000

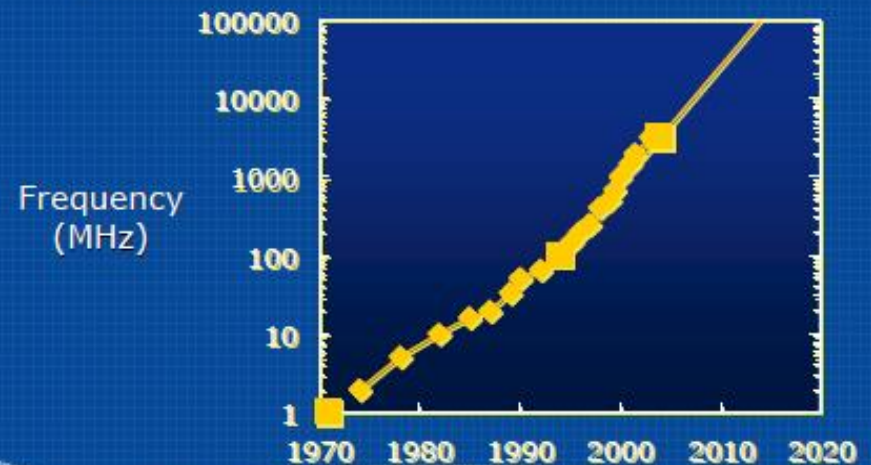


# Historical Driving Forces

## Shrinking Geometry



## Increased Frequency



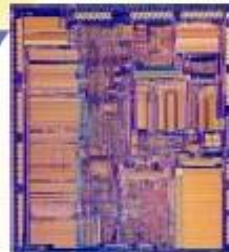
**1971**

4004 Processor  
2300 Transistors



**1978**

8008 Processor  
IBM PC



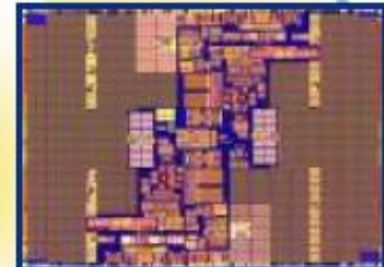
**1986**

1386 Processor  
32-bit



**1993**

Pentium Processor  
3.1M transistors



**2005**

Montecito  
1.7B Transistors

Courtesy, Intel

# The Act of Balancing

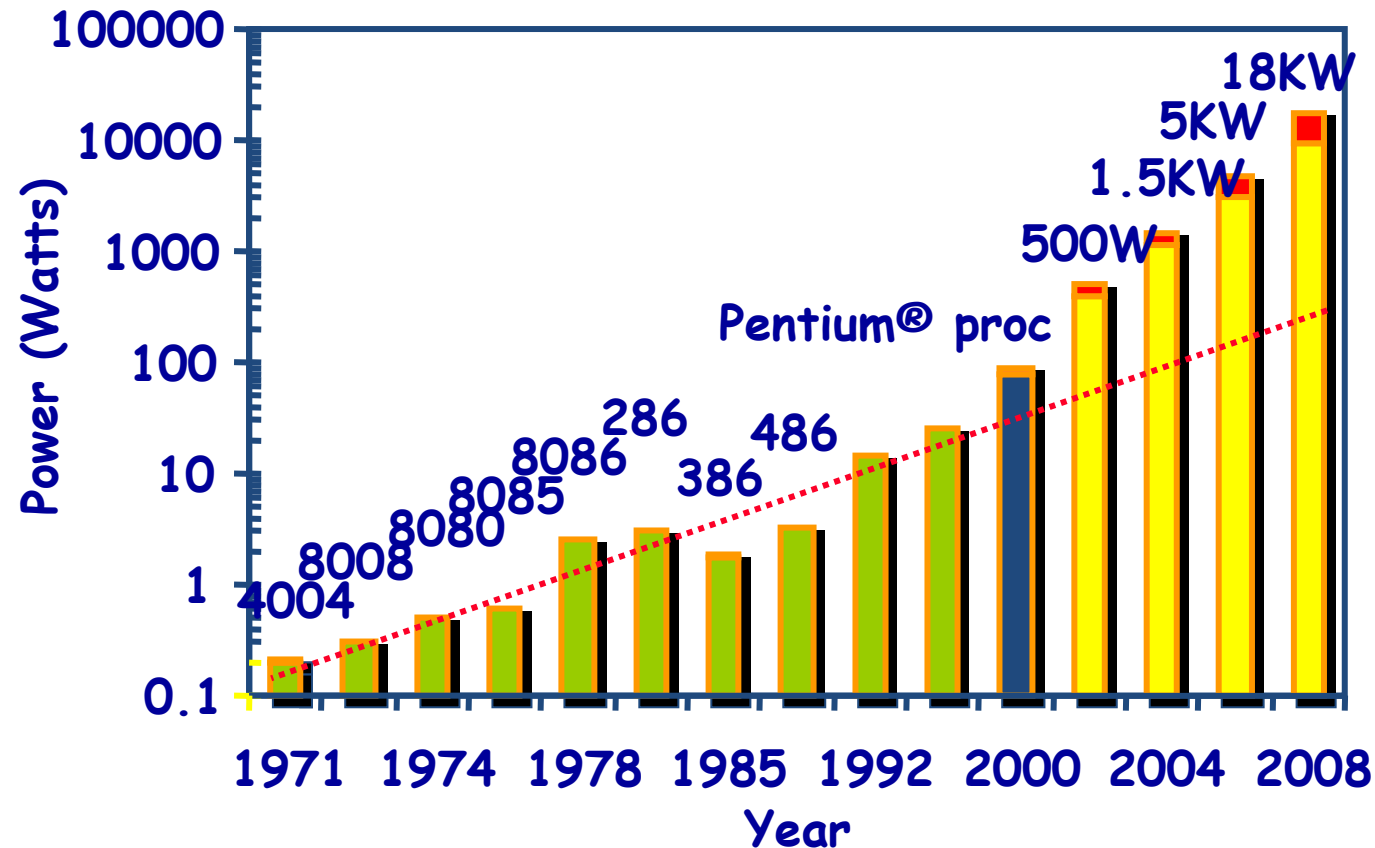


**Delivered Performance =  
Instructions Per Cycle (IPC) \* Frequency**

*Goal is higher performance and lower power*

**Power  $\propto$   $C_{\text{dynamic}}$  \* V \* V \* Frequency**

# Power will be a major problem



Power delivery and dissipation will be prohibitive



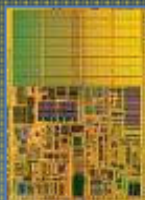
# Dual core with voltage scaling

## RULE OF THUMB

A 15%  
Reduction  
In Voltage  
Yields

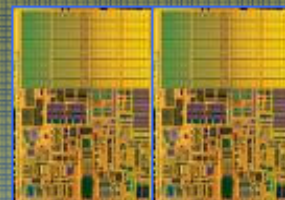
Frequency Reduction	Power Reduction	Performance Reduction
15%	45%	10%

### SINGLE CORE



Area = 1  
Voltage = 1  
Freq = 1  
Power = 1  
Perf = 1

### DUAL CORE



Area = 2  
Voltage = 0.85  
Freq = 0.85  
Power = 1  
Perf =  $\sim 1.8$

# ***A New Era...***

## ***THE OLD***

**Performance  
Equals Frequency**

**Unconstrained Power**

**Voltage Scaling**

## ***THE NEW***

**Performance  
Equals IPC**

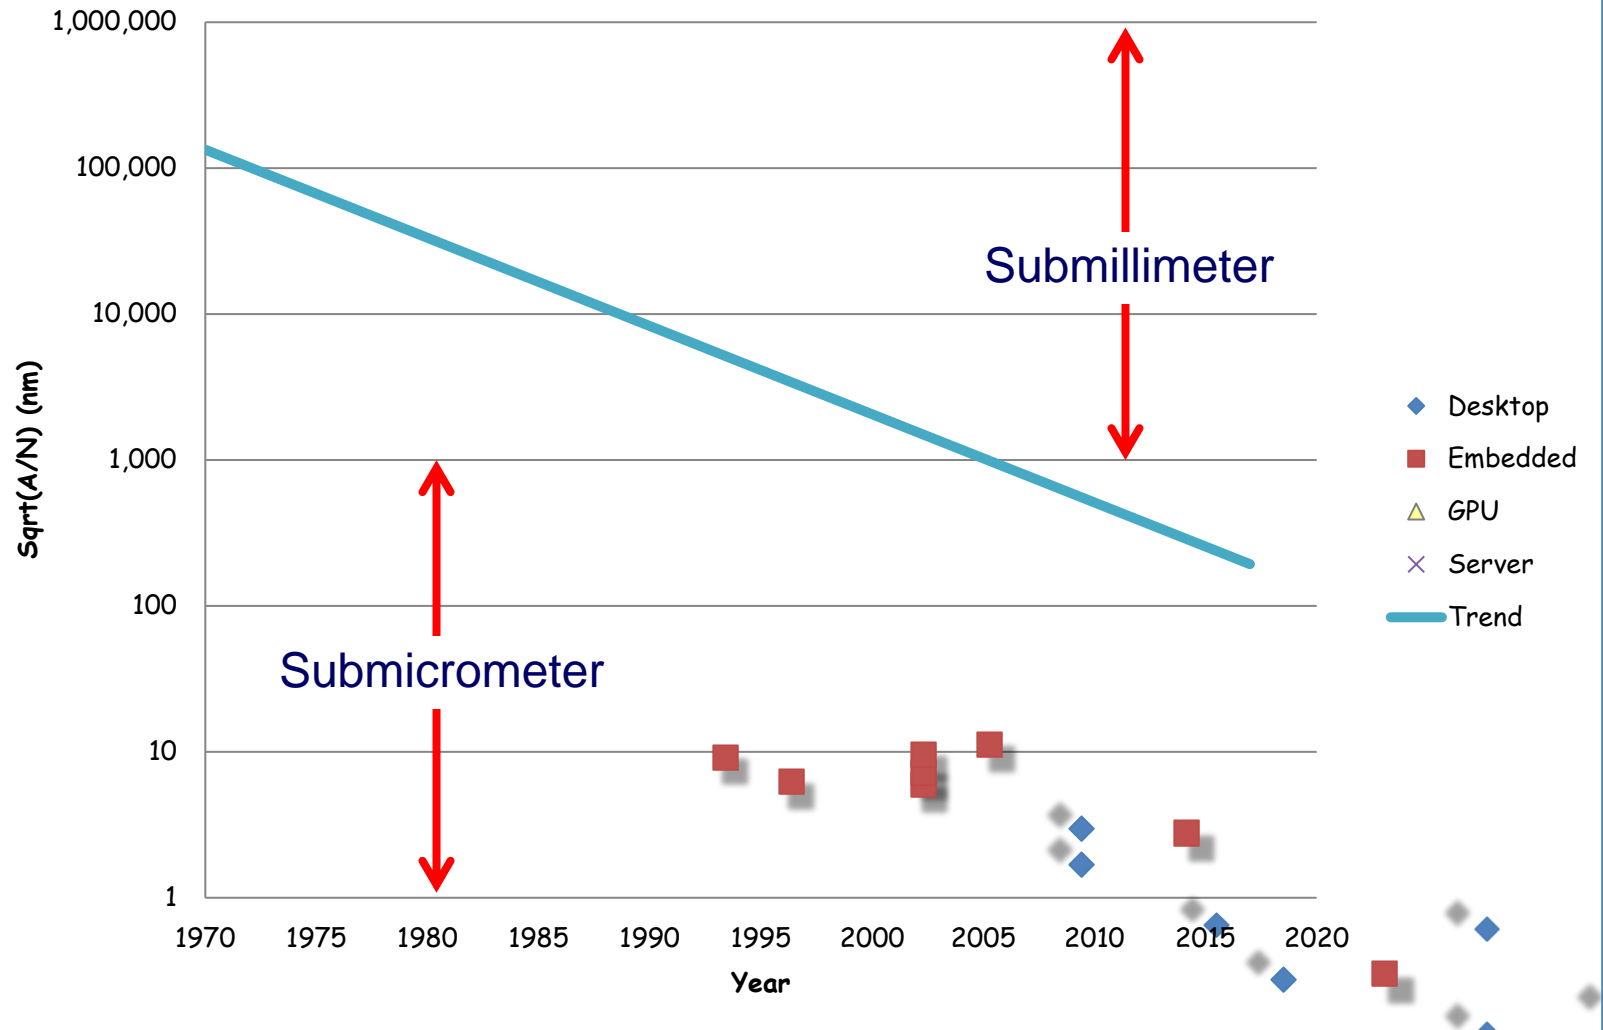
**Multi-Core**

**Power Efficiency**

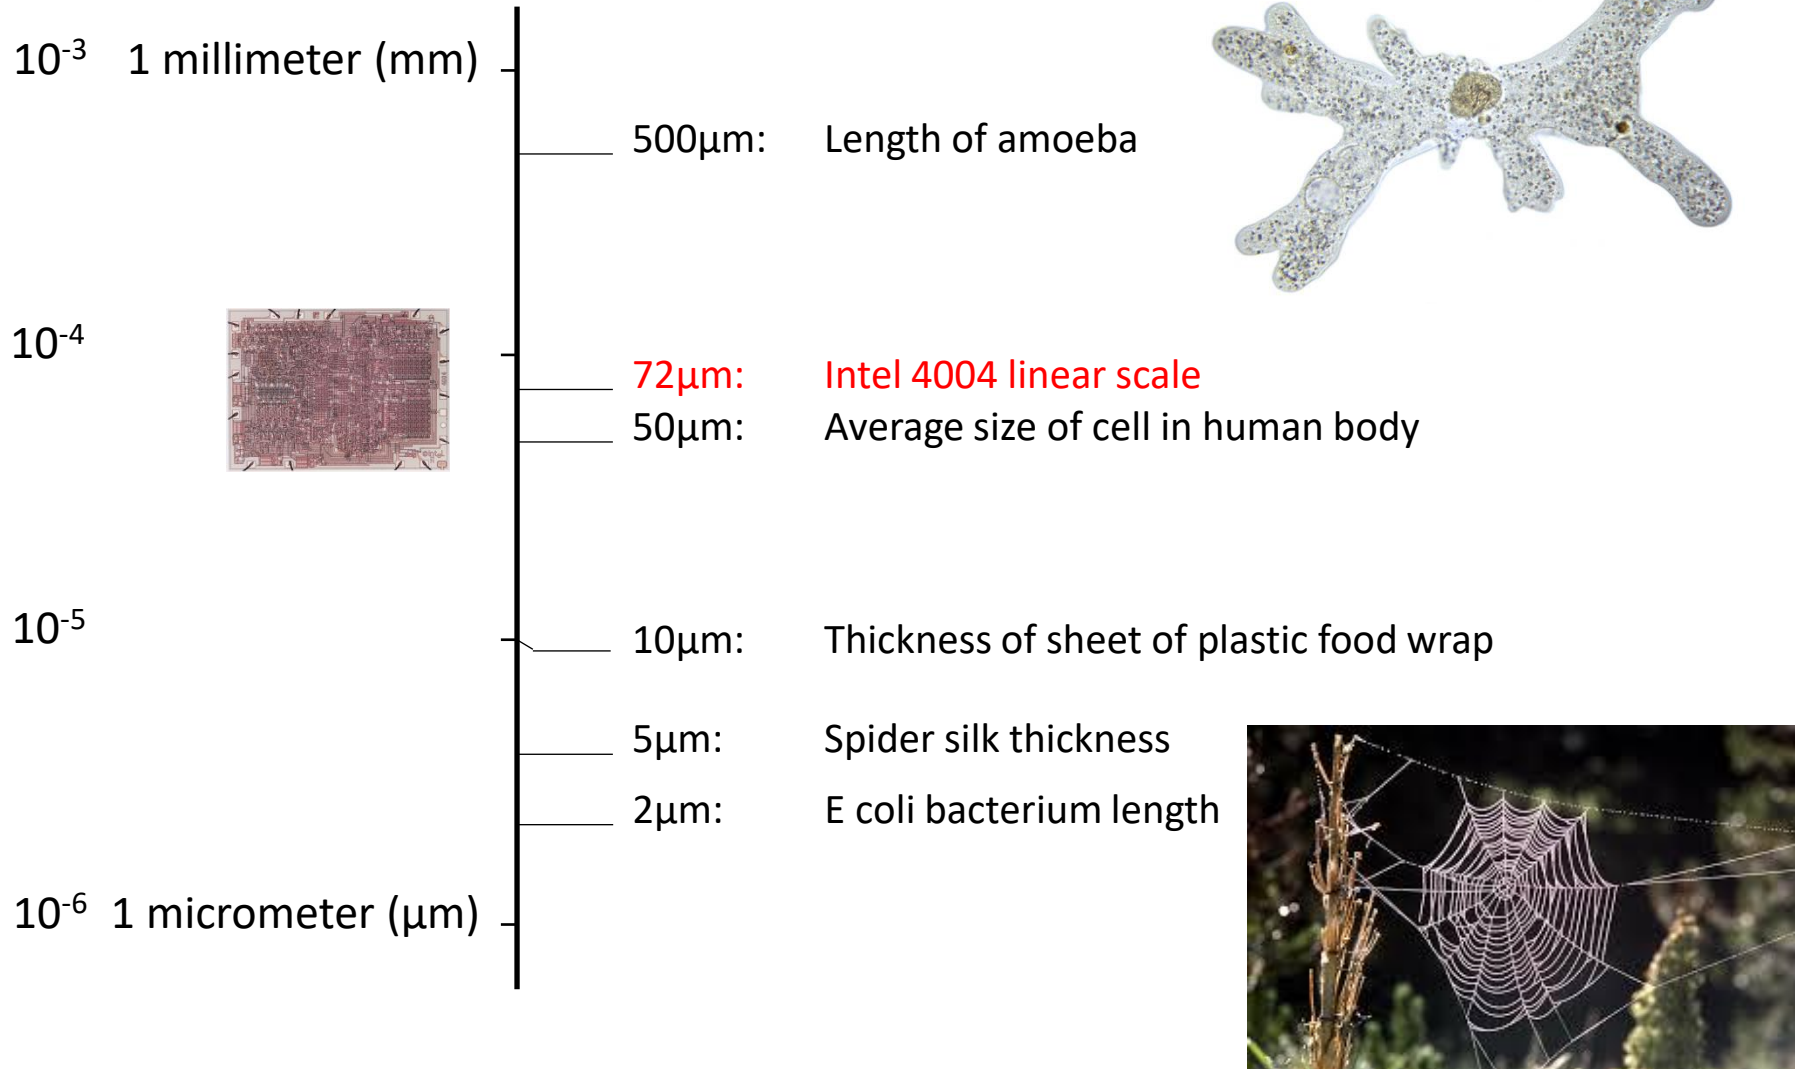
**Microarchitecture  
Advancements**

# Linear Scaling Trend

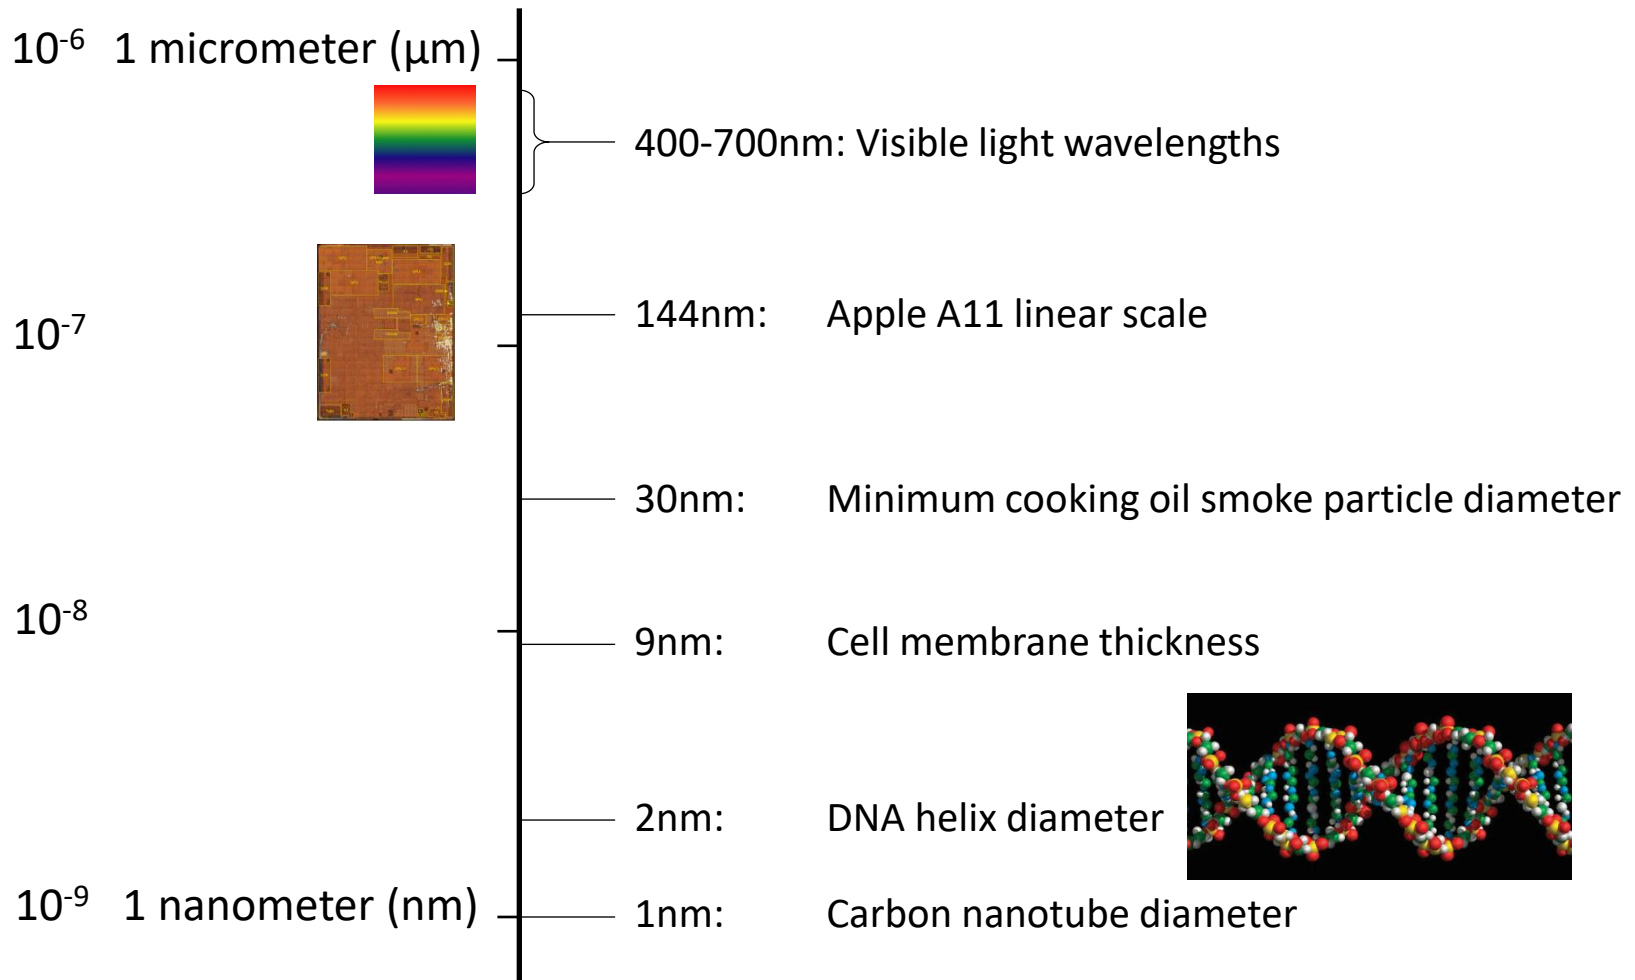
Linear Spacing by Year



# Submillimeter Dimensions

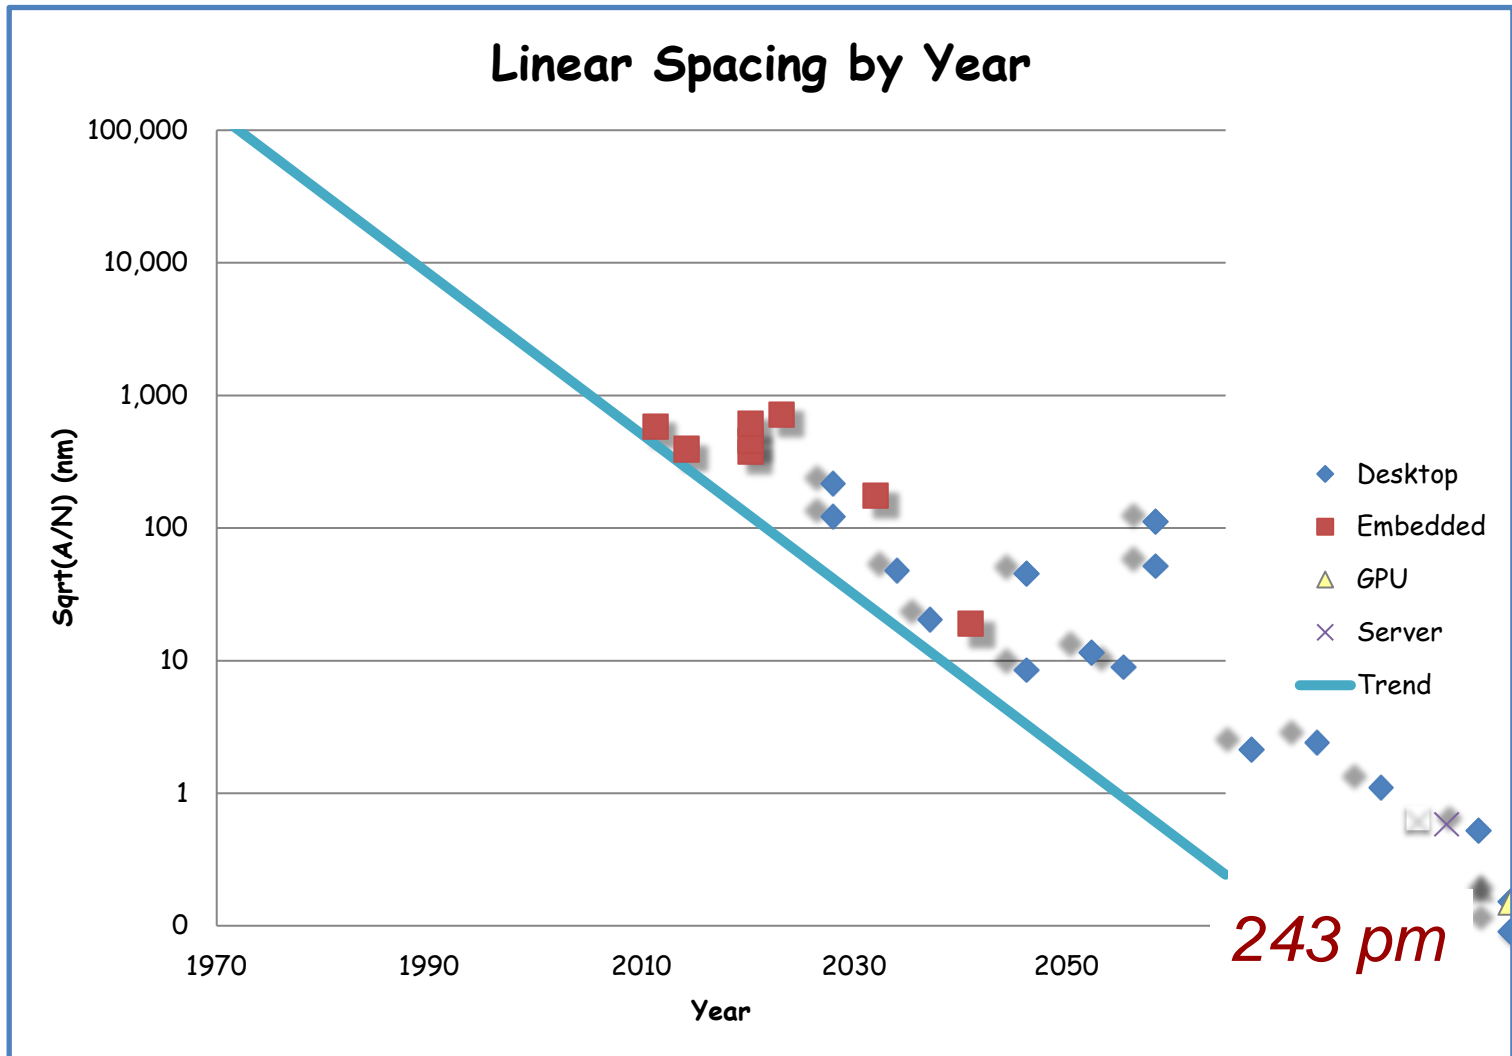


# Submicrometer Dimensions





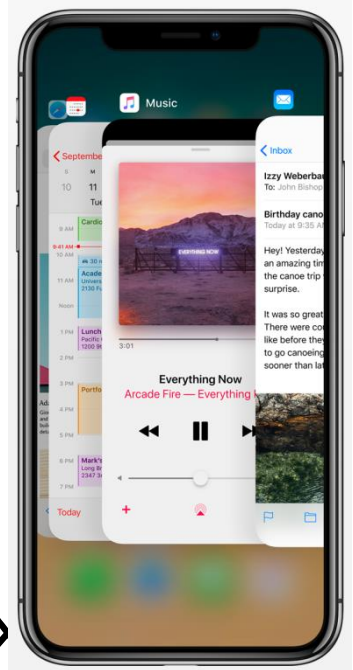
# Linear Scaling Extrapolation



# What Moore's Law Has Meant



- 1976 Cray 1
  - 250 M Ops/second
  - ~170,000 chips
  - 0.5B transistors
  - 5,000 kg, 115 KW
  - \$9M
  - 80 manufactured



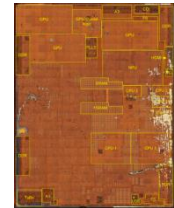
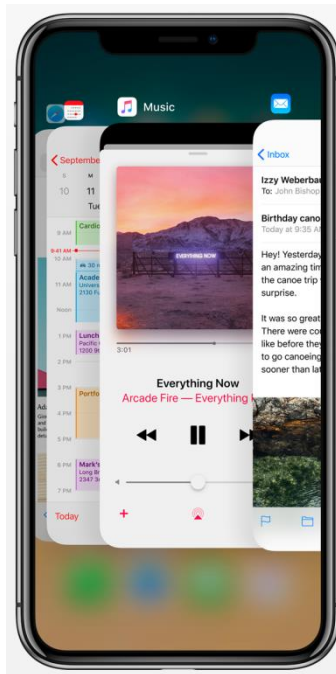
- 2017 iPhone >
  - > 10 B Ops/second
  - 16 chips
  - 4.3B transistors (CPU only)
  - 174 g, < 5 W
  - \$999
  - ~3 million sold in first 3 days

# What Moore's Law Has Meant

- 1965 Consumer Product



- 2017 Consumer Product

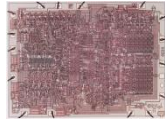


Apple A11 Processor  
4.3B transistors

# Visualizing Moore's Law to Date

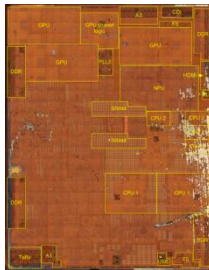
*If transistors were the size of a grain of sand*

Intel 4004  
1970  
2,300 transistors



0.1 g

Apple A11  
2017  
4.3 B transistors



189 kg

# What Moore's Law Has Meant

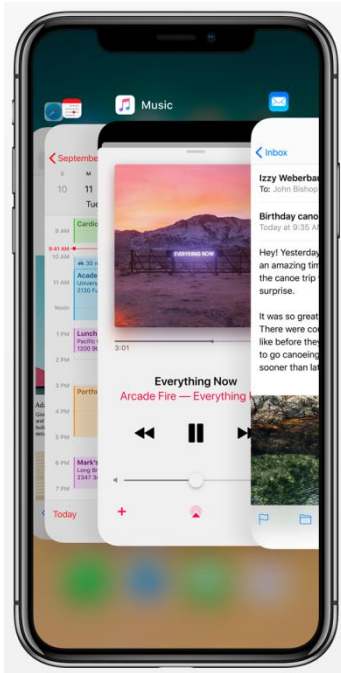
iPhone	Released with	Release date
iPhone (1st Gen.)	iPhone OS 1.0	June 29, 2007
iPhone 3G	iPhone OS 2.0	July 11, 2008
iPhone 3GS	iPhone OS 3.0	June 19, 2009
iPhone 4	iOS 4.0	June 21, 2010
iPhone 4S	iOS 5.0	October 14, 2011
iPhone 5	iOS 6.0	September 21, 2012
iPhone 5C	iOS 7.0	September 20, 2013
iPhone 5S	iOS 7.0	September 20, 2013
iPhone 6 (Plus)	iOS 8.0	September 19, 2014
iPhone 6S (Plus)	iOS 9.0	September 25, 2015
iPhone SE	iOS 9.3	March 31, 2016
iPhone 7 (Plus)	iOS 10.0	September 16, 2016
iPhone 8 (Plus)	iOS 11.0	September 22, 2017
iPhone X	iOS 11.0.1	November 3, 2017

12 generations of iPhone since 2007



# What Moore's Law Could Mean

- 2017 Consumer Product
- 2065 Consumer Product



- Portable
- Low power
- Will drive markets & innovation

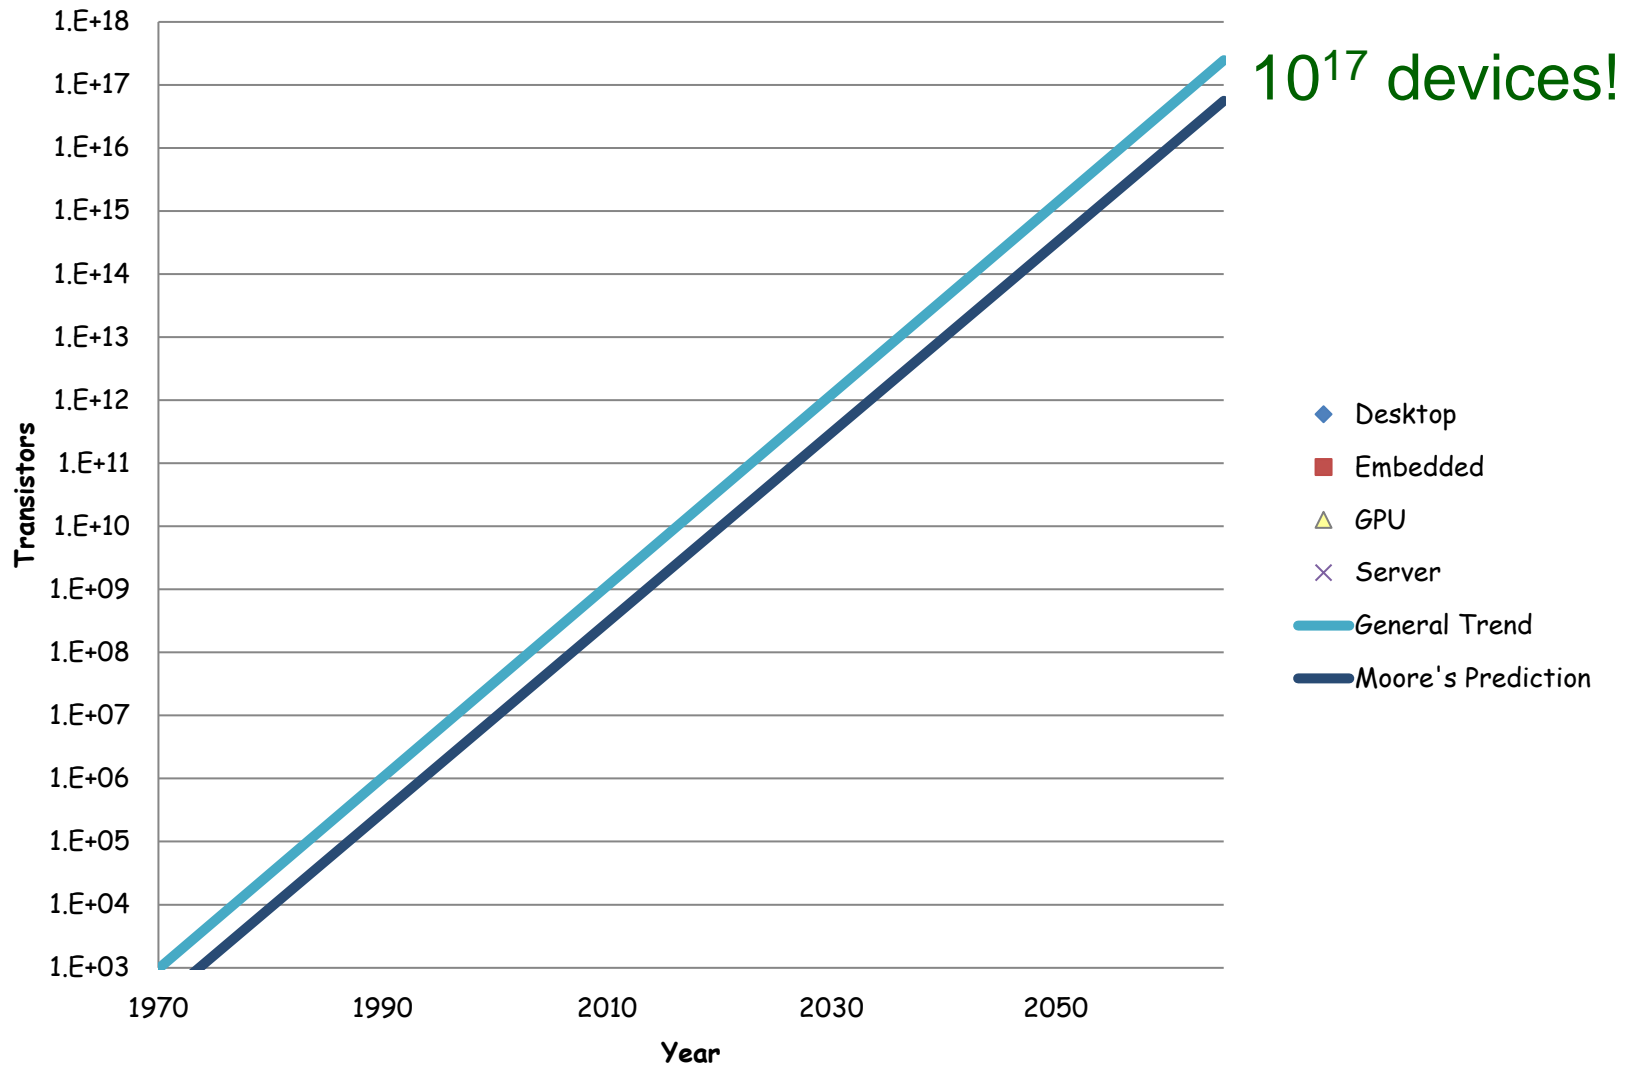
# Requirements for Future Technology

- Must be suitable for portable, low-power operation
  - Consumer products
  - Internet of Things components
- Must be inexpensive to manufacture
  - Comparable to current semiconductor technology
    - $O(1)$  cost to make chip with  $O(N)$  devices
- Need not be based on transistors
  - Memristors, carbon nanotubes, DNA transcription, ...
  - Possibly new models of computation
  - But, still want lots of devices in an integrated system



# Moore's Law: 100 Years

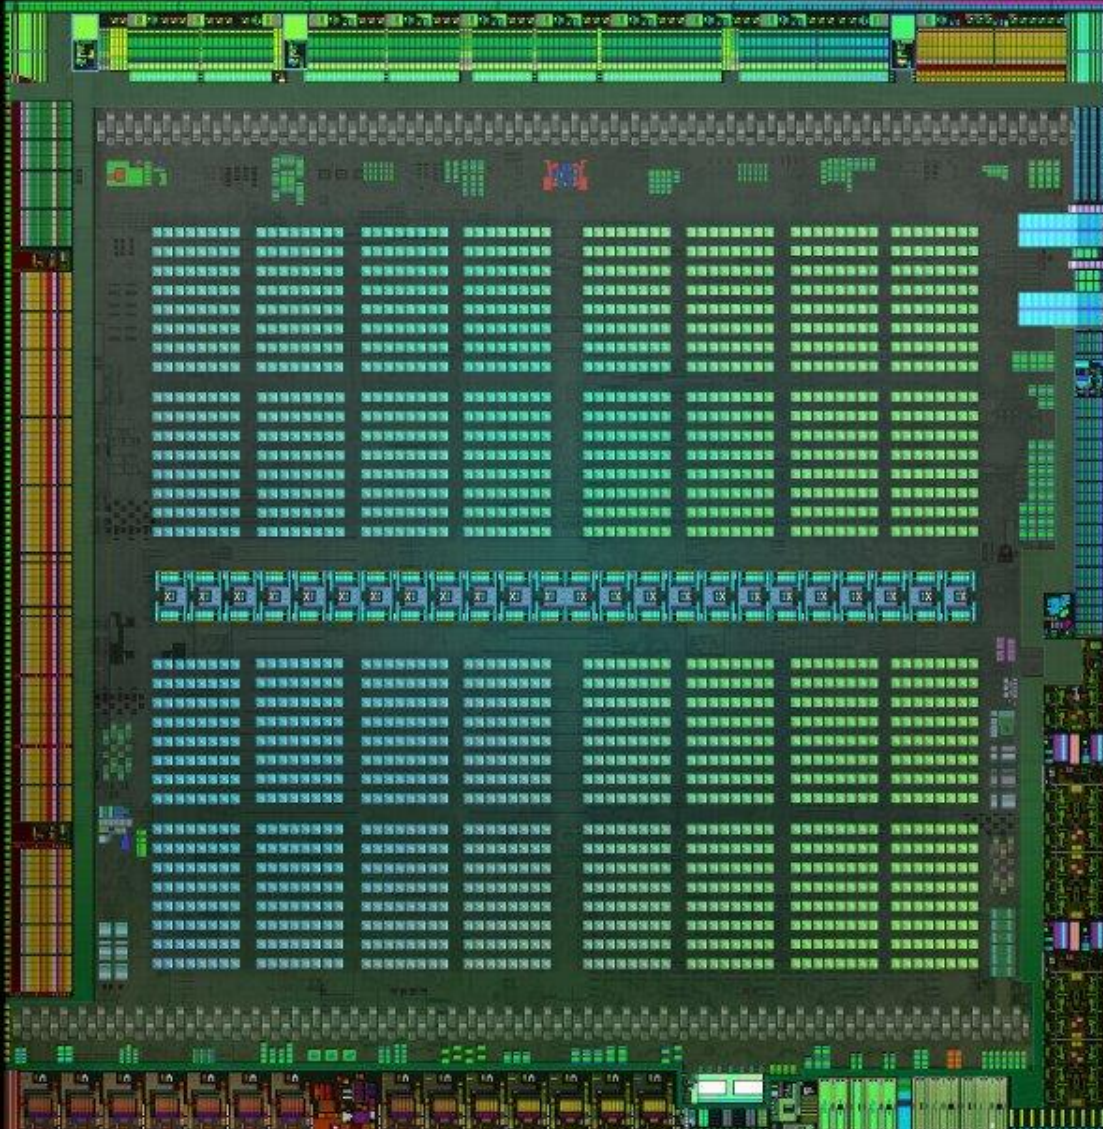
## Device Count by Year



# NVIDIA GPU Specification Comparison

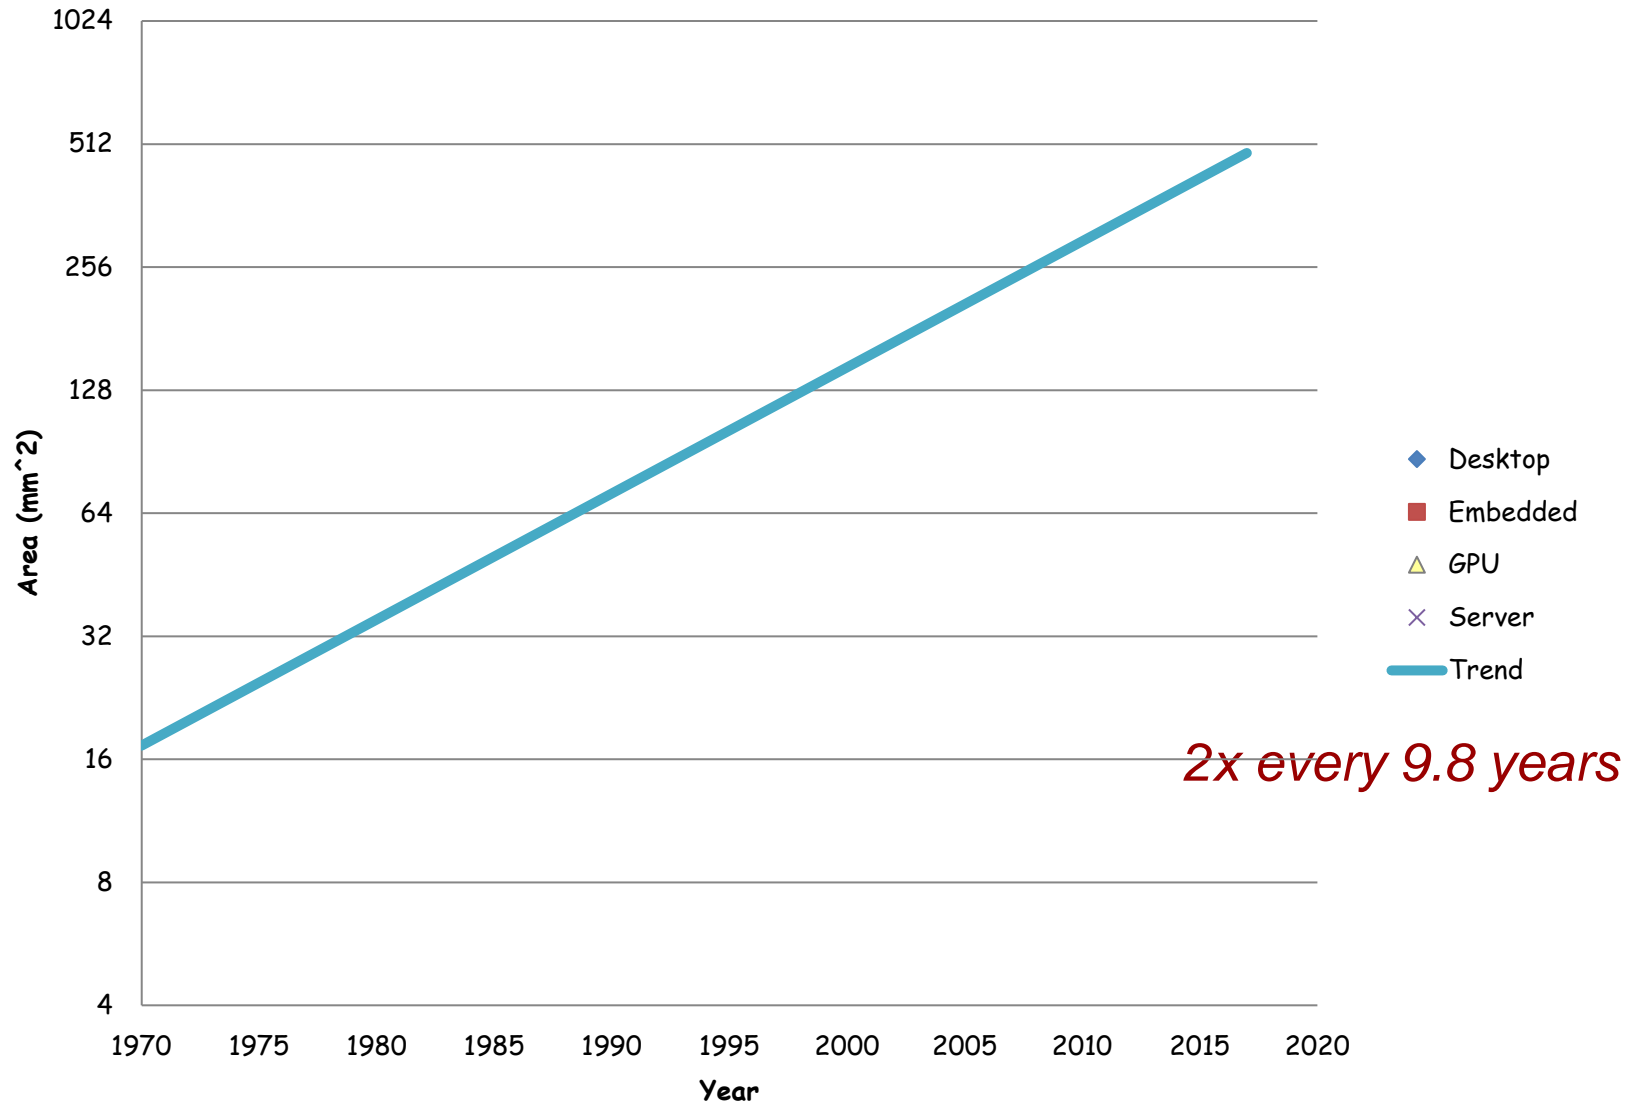
NVIDIA GV100 Volta  
2017  
21.1 B transistors  
815 mm<sup>2</sup>

	GV100	GP100	GV100
CUDA Cores	5376	3840	2880
Tensor Cores	672		
SMs	84		
CUDA Cores/SM	64		
Tensor Cores/SM	8		
Texture Units	336		
Memory	HBM		
Memory Bus Width	4096		
Shared Memory	128KB, Con		
L2 Cache	6MB		
Half Precision	2:1 (Ve		
Double Precision	1:2		
Die Size	815mm		
Transistor Count	21.1		
TDP	300W		
Manufacturing Process	TSMC 12n		
Architecture	Volta		



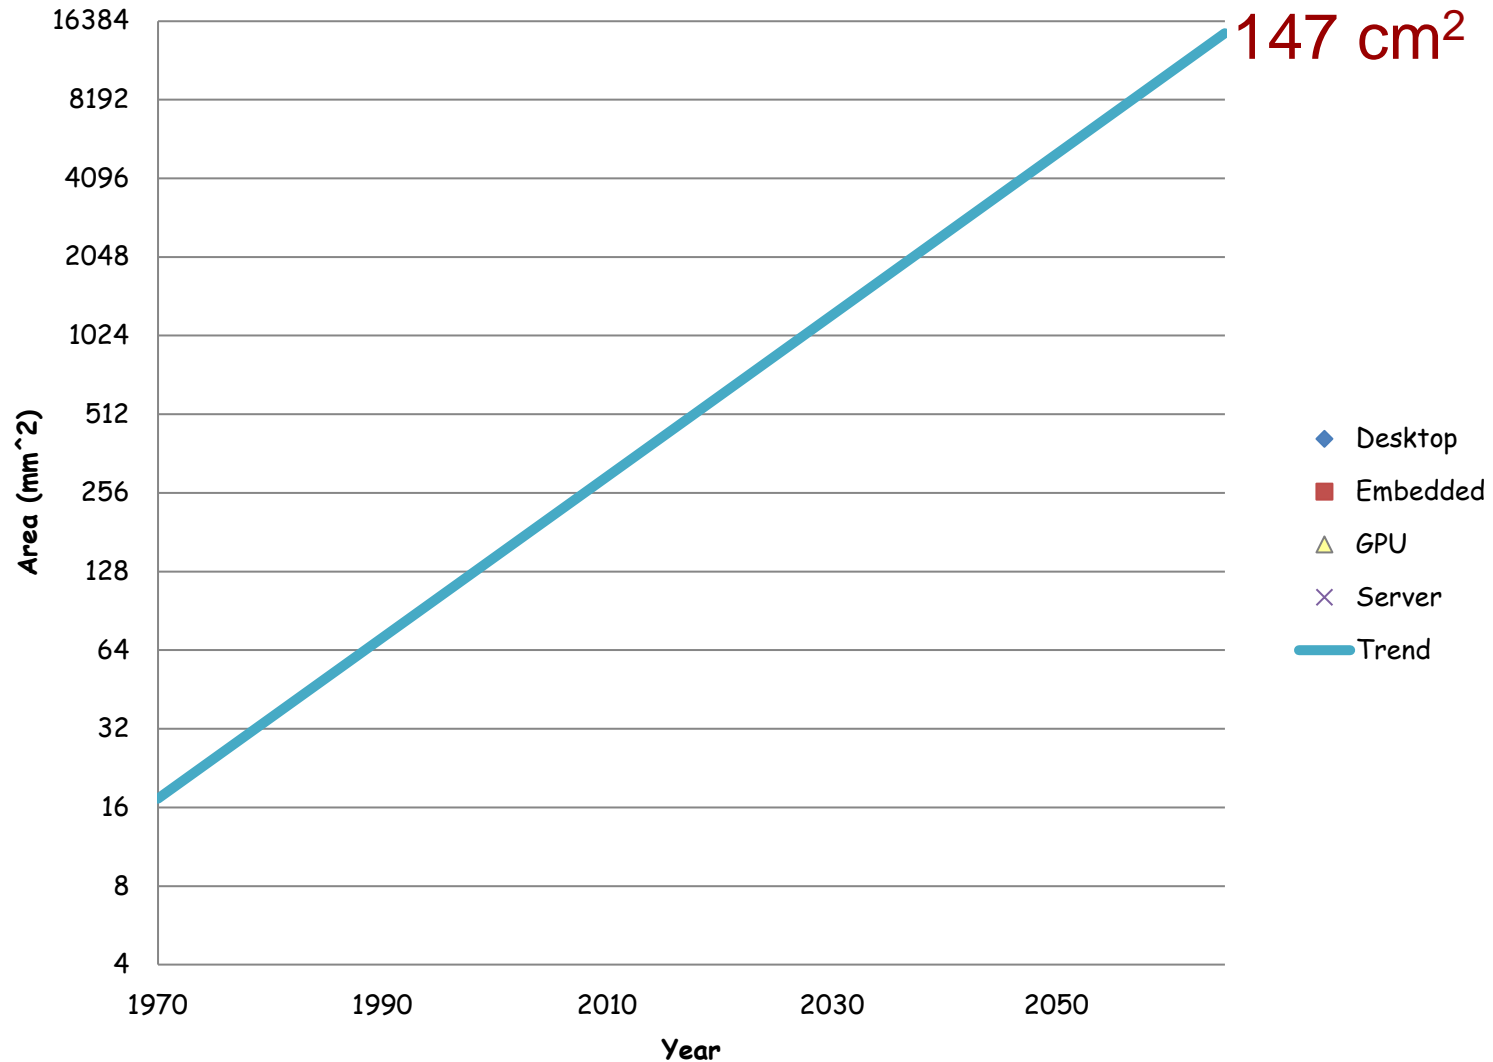
# Chip Size Trend

Area by Year



# Chip Size Extrapolation

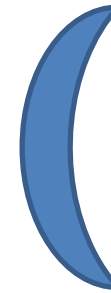
Area by Year



# This Course

- Micro-architecture: how to implement an architecture in hardware
- Processor:
  - Datapath: functional blocks
  - Control: control signals

This Course



Application Software	programs
Operating Systems	device drivers
Architecture	instructions registers
Micro-architecture	datapaths controllers
Logic	adders memories
Digital Circuits	AND gates NOT gates
Analog Circuits	amplifiers filters
Devices	transistors diodes
Physics	electrons



# What is Computer Architecture?

- The science and art of designing, selecting, and interconnecting hardware components and designing the hardware/software interface to create a computing system that meets functional, performance, energy consumption, cost, and other specific goals.
- We will soon distinguish between the terms *architecture*, and *microarchitecture*.

# Why Study Computer Architecture?

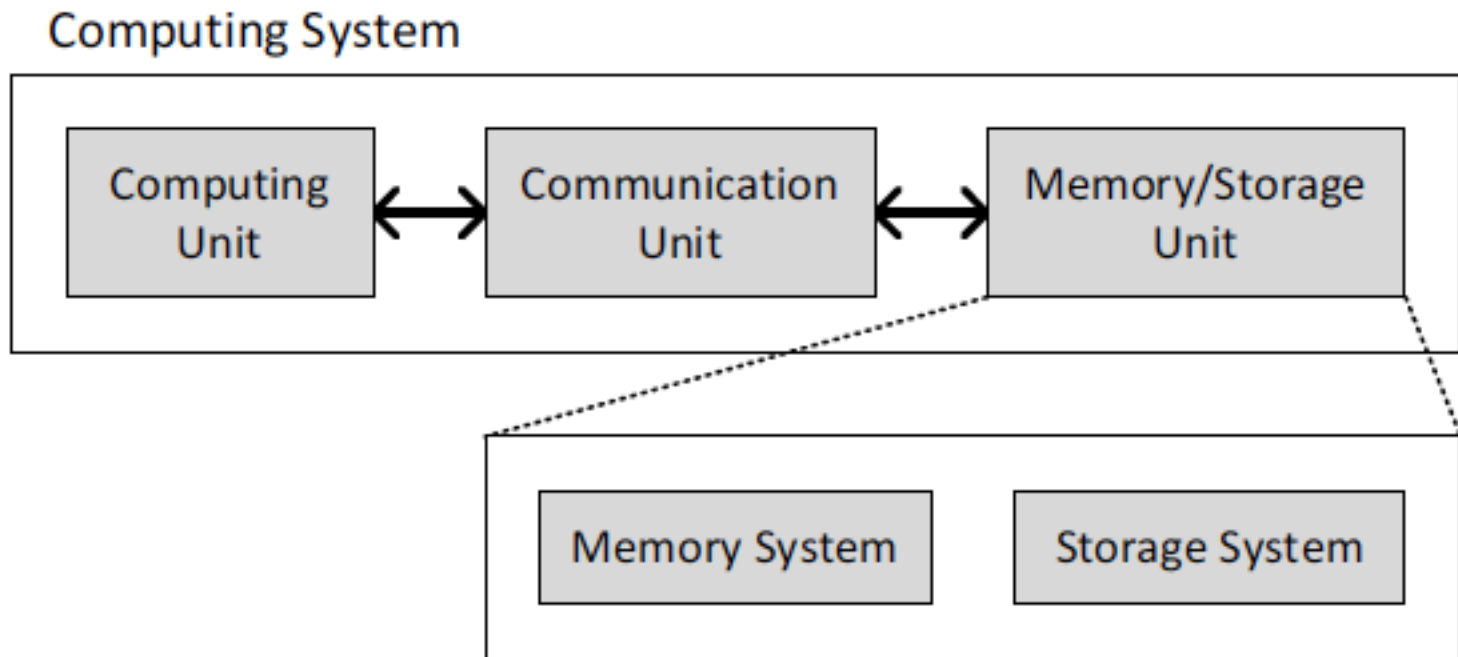
- **Enable better systems:** make computers faster, cheaper, smaller, more reliable, ...
  - By exploiting advances and changes in underlying technology/circuits
- **Enable new applications**
  - Life-like 3D visualization 20 years ago?
  - Virtual reality?
- **Enable better solutions to problems**
  - Software innovation is built into trends and changes in computer architecture
    - > 50% performance improvement per year has enabled
- **Understand why computers work the way they do**

# Computer Architecture Today

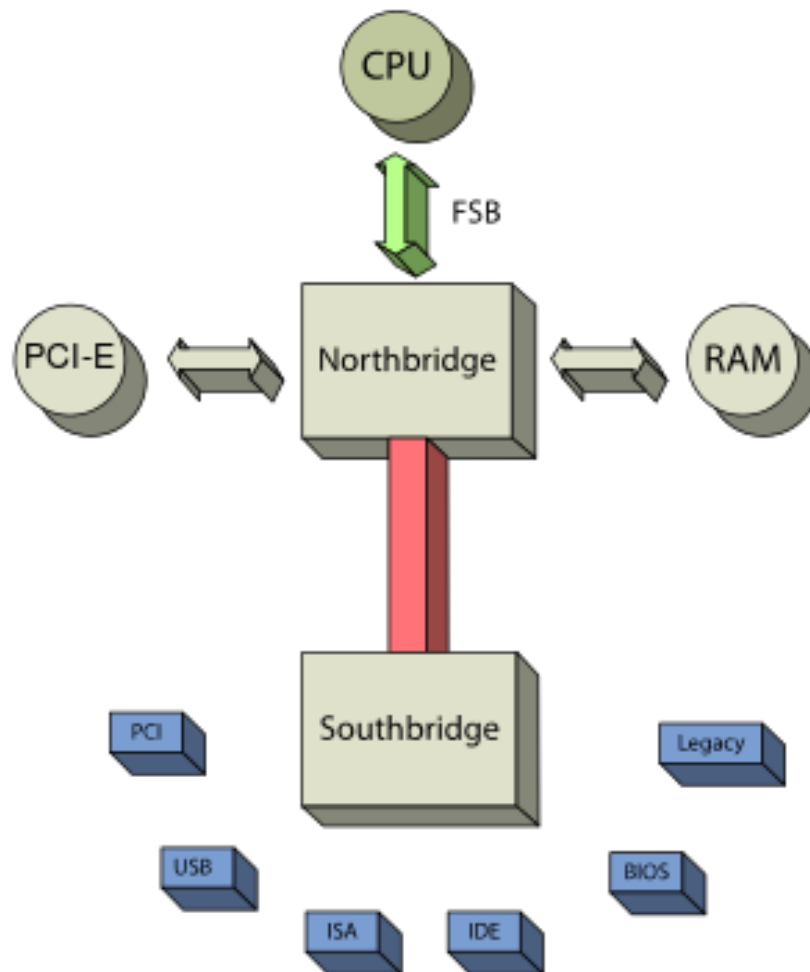
- Today is a very exciting time to study computer architecture
- Industry is in a large paradigm shift (to multi-core)
- Many problems motivating and caused by the shift
  - Power/energy constraints
  - Complexity of design → multi-core
  - Technology scaling → new technologies
  - Memory wall/gap
  - Reliability wall/issues
  - Programmability wall/problem

# What is A Computer?

- Three key components
- Computation
- Communication
- Storage (memory)



# Typical PC Architecture: **northbridge** or **host bridge**



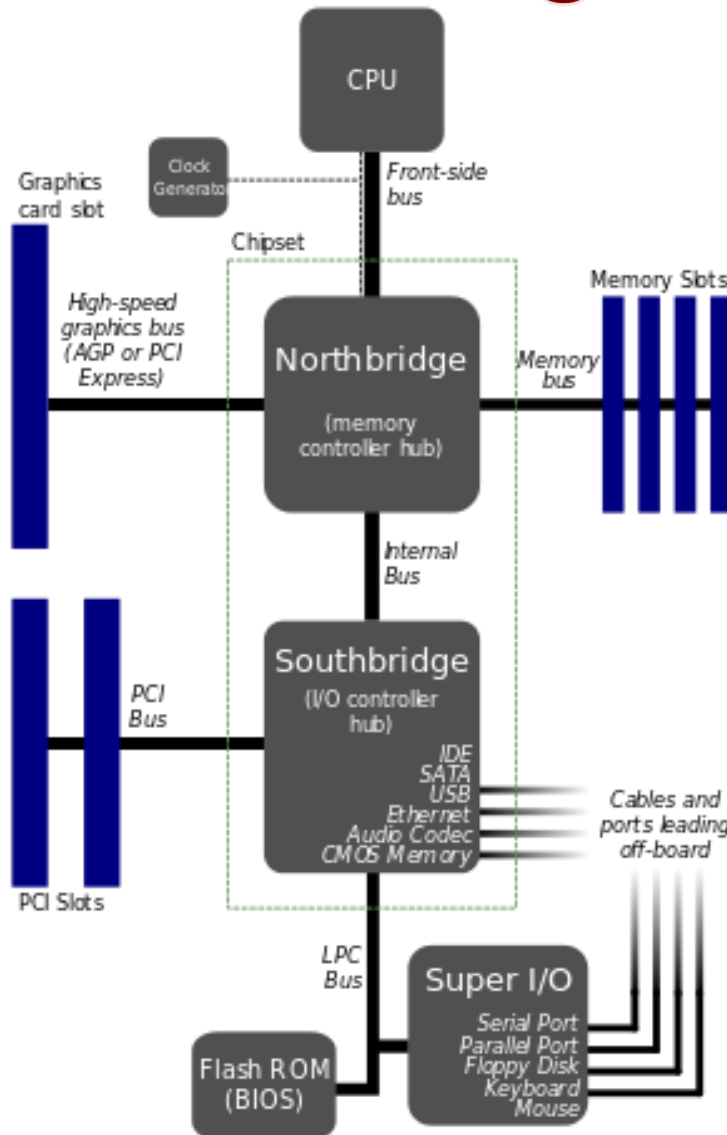
front-side bus (FSB)

I/O controller hub



# (Typical PC Architecture)

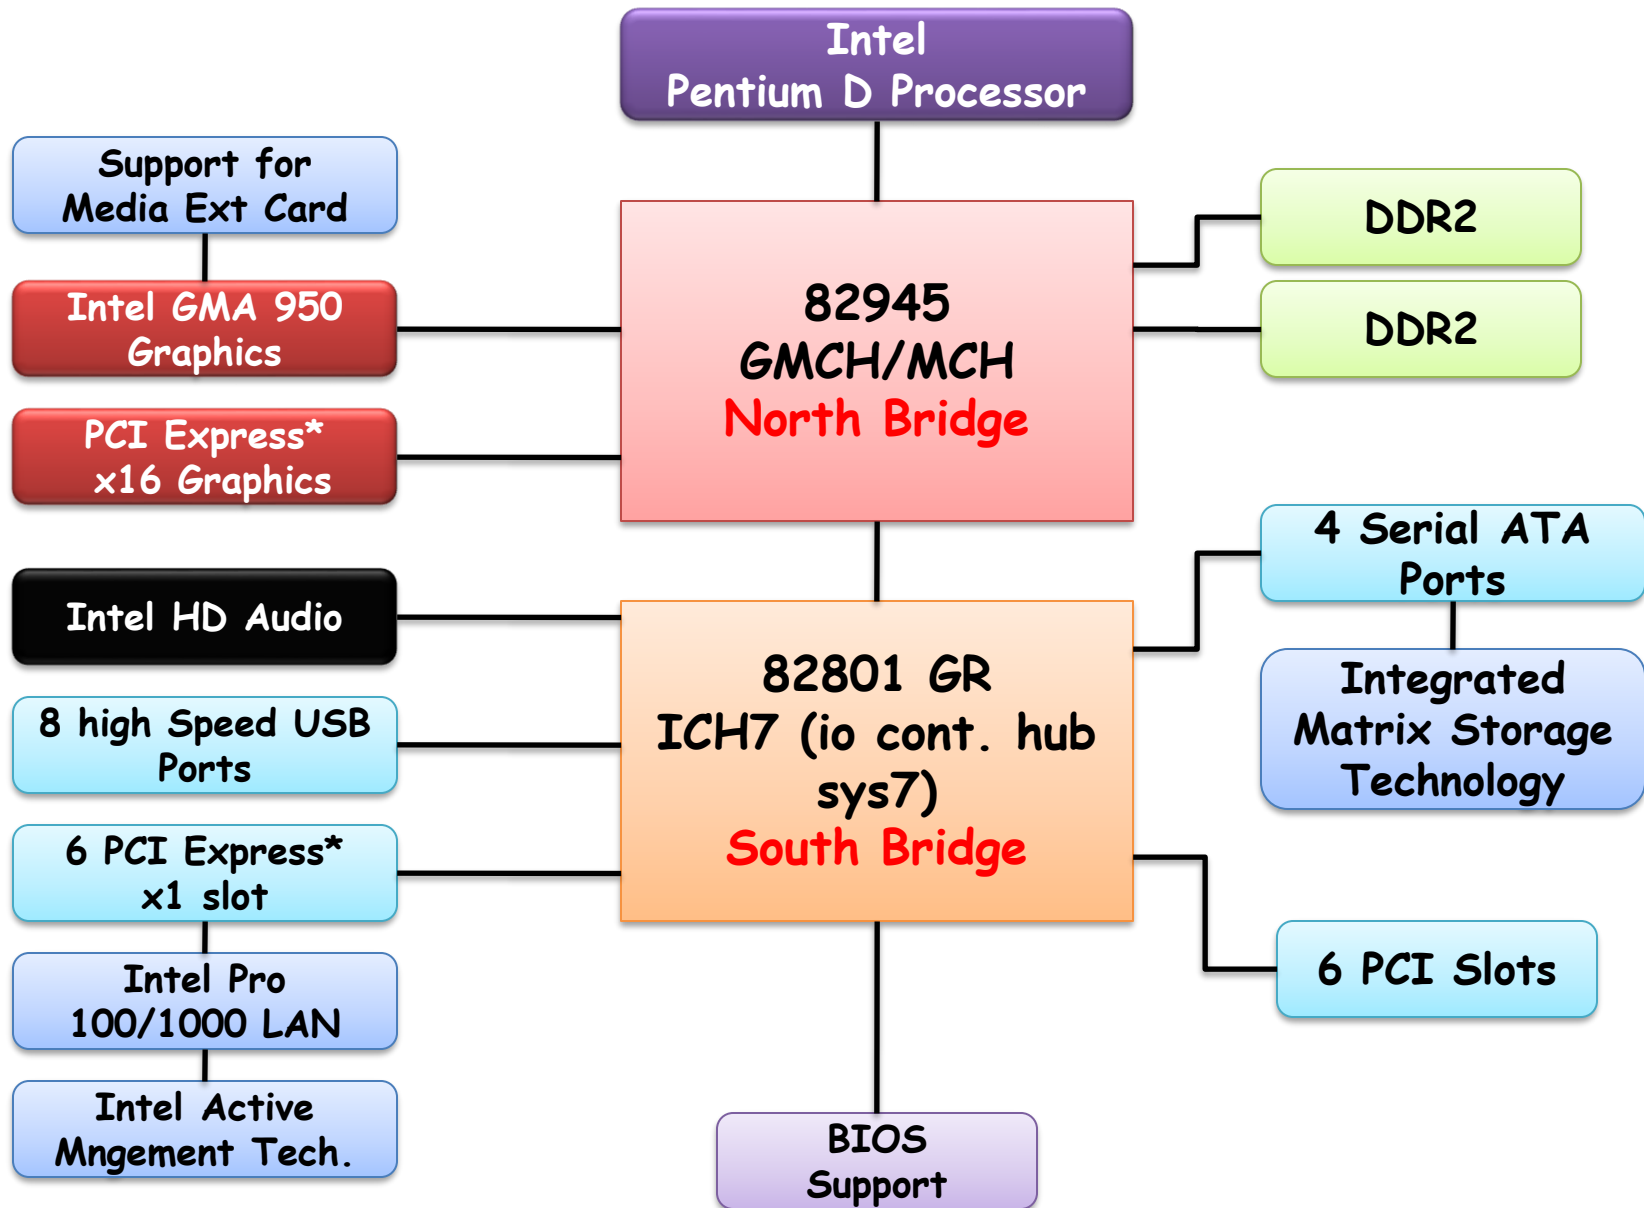
## northbridge or host bridge



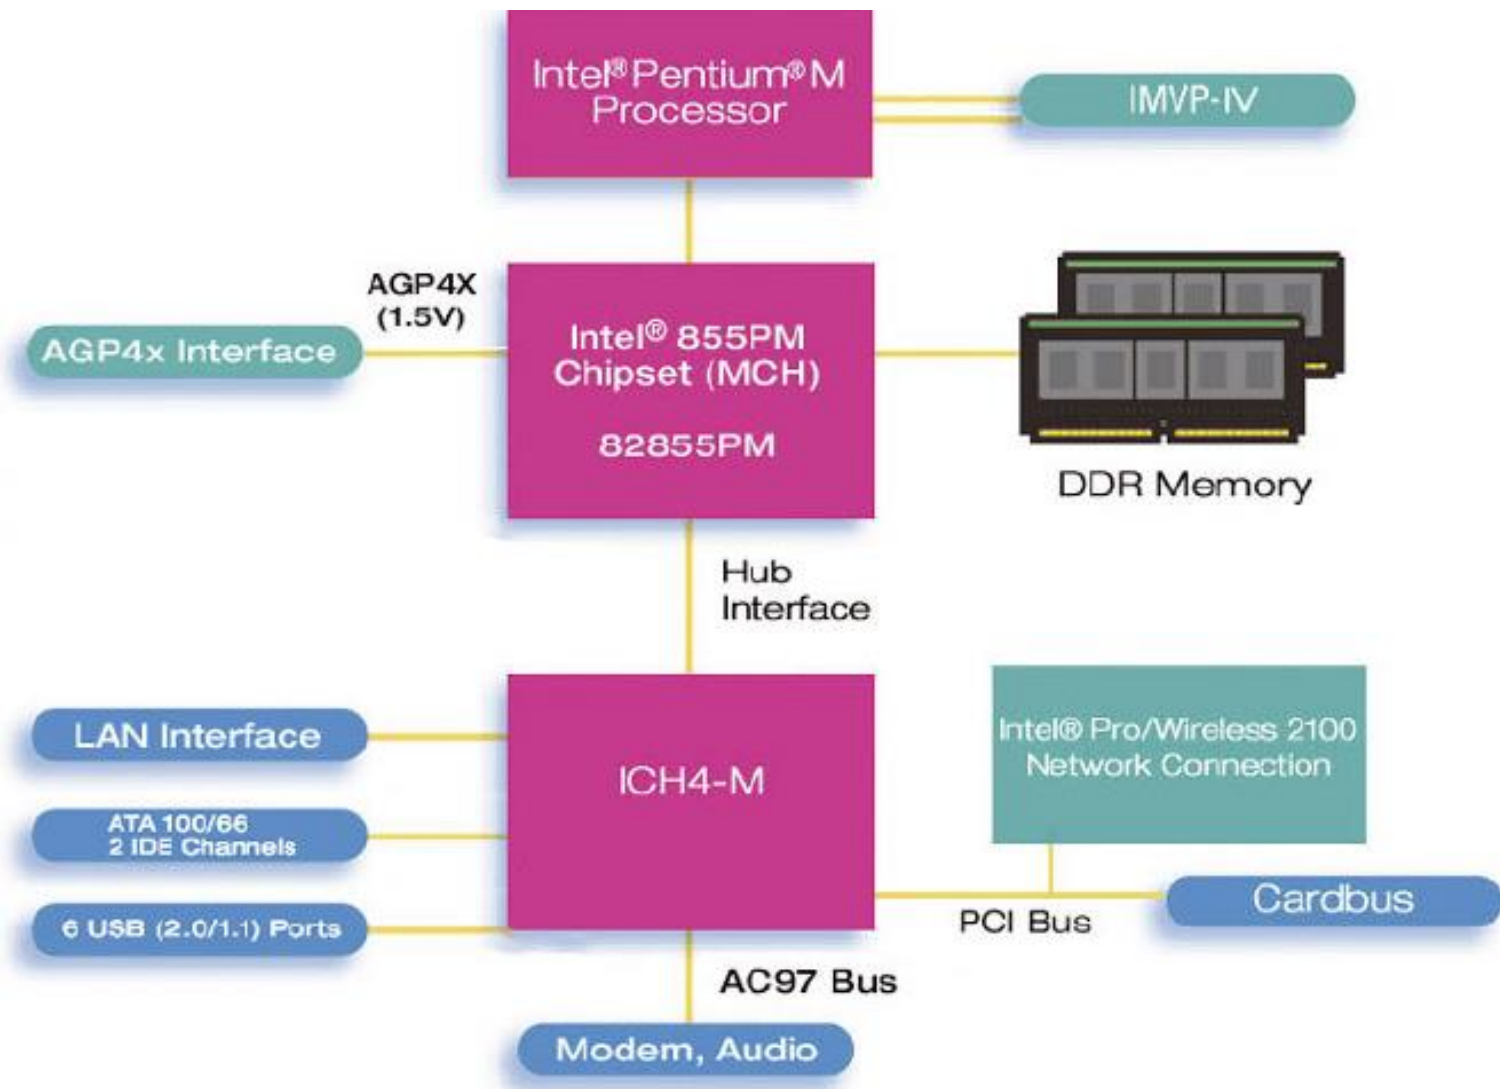
**Graphics and Memory Controller Hub (GMCH)**

**I/O Controller Hub (ICH)**

# Intel 945 Express Chipset



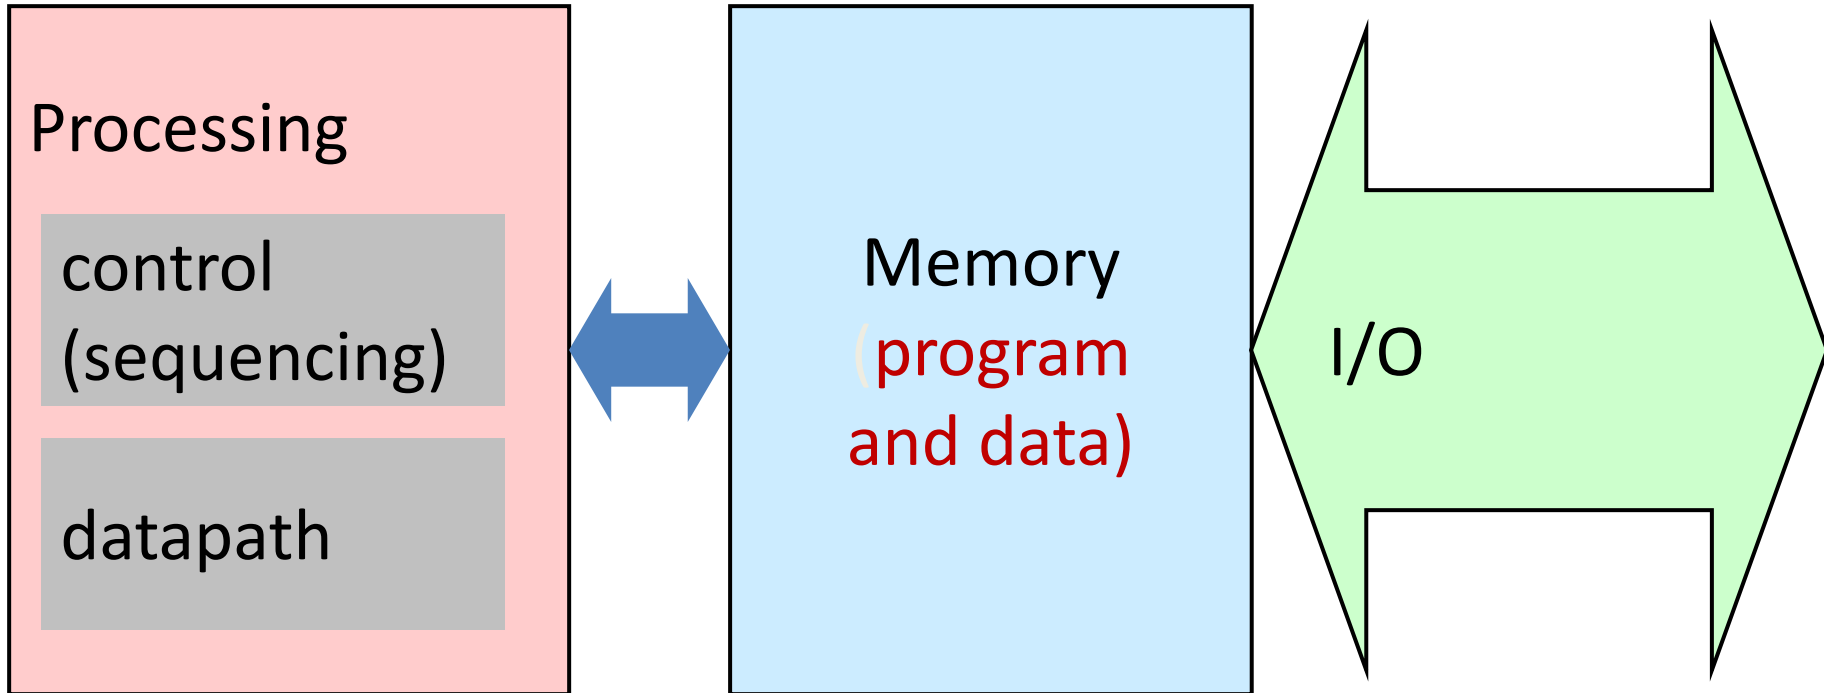
# What is A Computer?



Courtesy, Intel

# What is A Computer?

- We will cover all three components



# The Von Neumann Model/Architecture

- Also called *stored program computer* (instructions in memory). Two key properties:
  - Stored program
    - Instructions stored in a linear memory array
    - Memory is unified between instructions and data
      - The interpretation of a stored value depends on the control signals
- When is a value interpreted as an instruction?
- Sequential instruction processing
    - One instruction processed (fetched, executed, and completed) at a time
    - Program counter (instruction pointer) identifies the current instr.
    - Program counter is advanced sequentially except for control transfer instructions



# Dataflow Model (of a Computer)

- Von Neumann model: An instruction is fetched and executed in **control flow order**
  - As specified by the **instruction pointer**
  - Sequential unless explicit control flow instruction
- Dataflow model: An instruction is fetched and executed in **data flow order**
  - i.e., when its operands are ready
  - i.e., there is **no instruction pointer**
  - Instruction ordering specified by data flow dependence
    - Each instruction specifies “who” should receive the result
    - An instruction can “fire” whenever all operands are received
  - Potentially many instructions can execute at the same time
    - Inherently more parallel

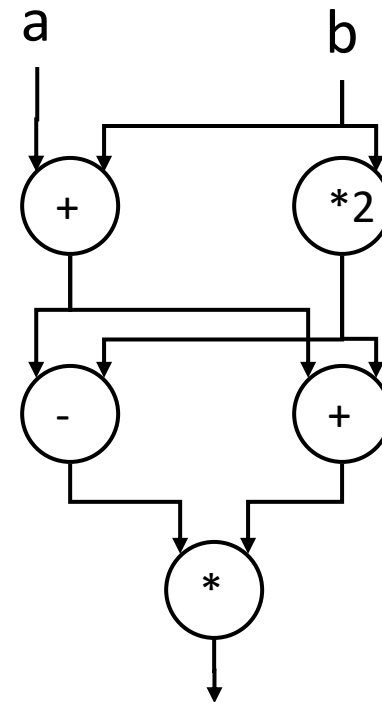
# Von Neumann vs Dataflow

## ■ Consider a von Neumann program

- What is the significance of the program order?
- What is the significance of the storage locations?

```
v <= a + b;  
w <= b * 2;  
x <= v - w  
y <= v + w  
z <= x * y
```

Sequential



Dataflow

## ■ Which model is more natural to you as a programmer?

# ISA vs. Microarchitecture Level Tradeoff

- A similar tradeoff (control vs. data-driven execution) can be made at the microarchitecture level
- ISA: Specifies how the programmer sees instructions to be executed
  - Programmer sees a sequential, control-flow execution order vs.
  - Programmer sees a data-flow execution order
- Microarchitecture: How the underlying implementation actually executes instructions
  - Microarchitecture can execute instructions in any order as long as it obeys the semantics specified by the ISA when making the instruction results visible to software
    - Programmer should see the order specified by the ISA

# ISA vs. Microarchitecture Level

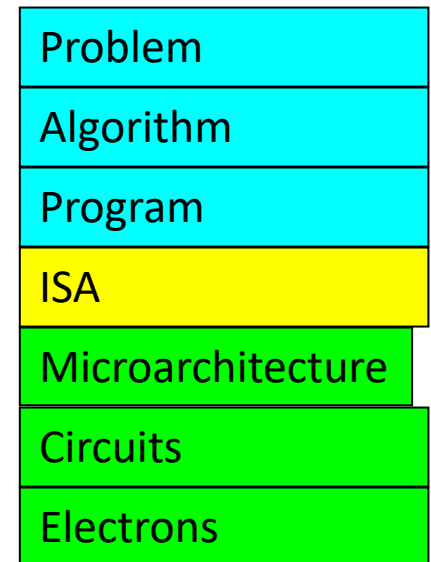
- ISA: Specifies how the programmer sees instructions to be executed
- Microarchitecture: How the underlying implementation actually executes instructions
  - Microarchitecture can execute instructions in any order as long as it obeys the semantics specified by the ISA when making the instruction results visible to software
    - Programmer should see the order specified by the ISA

# The Von-Neumann Model

- All major *instruction set architectures* today use this model
  - x86, ARM, MIPS, SPARC, Alpha, POWER
- Underneath (at the microarchitecture level), the execution model of almost all *implementations (or, microarchitectures)* is very different
  - Pipelined instruction execution: *Intel 80486 uarch*
  - Multiple instructions at a time: *Intel Pentium uarch*
  - Out-of-order execution: *Intel Pentium Pro uarch*
  - Separate instruction and data caches
- But, what happens underneath that is *not* consistent with the von Neumann model is *not* exposed to software
  - Difference between ISA and microarchitecture

# ISA vs. Microarchitecture

- ISA
  - Agreed upon interface between software and hardware
    - SW/compiler assumes, HW promises
  - What the software writer needs to know to write and debug system/user programs
- Microarchitecture
  - Specific implementation of an ISA
  - Not visible to the software
- Microprocessor
  - **ISA, uarch**, circuits
  - “Architecture” = ISA + microarchitecture





# CISC versus RISC

## CISC

Emphasis on hardware

Includes multi-clock  
complex instructions

Memory-to-memory:  
"LOAD" and "STORE"  
incorporated in instructions

Small code sizes,  
high cycles per second

Transistors used for storing  
complex instructions

## RISC

Emphasis on software

Single-clock,  
reduced instruction only

Register to register:  
"LOAD" and "STORE"  
are independent instructions

Low cycles per second,  
large code sizes

Spends more transistors  
on memory registers