



CS354: DATABASE

Relational Algebra, Tuple and Domain Relational
Calculus

1

QUERY LANGUAGES

- Language in which user requests information from the database
- Categories of languages
 - Procedural:
 - Relational algebra
 - Non-procedural, or declarative:
 - Tuple relational calculus, Domain relational calculus
- These languages form underlying basis of SQL query languages

RELATIONAL ALGEBRA

- Procedural language
- Six basic operators
 - Selection: σ
 - Projection: Π
 - Union: \cup
 - Set difference: $-$
 - Cartesian product: \times
 - Rename: ρ
- The operators take one or two relations as inputs and produce a new relation as a result

SELECTION OPERATION

- Relation r

A	B	C	D
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

$$\sigma_{(A=B) \wedge (D>5)}(r)$$

A	B	C	D
α	α	1	7
β	β	23	10

PROJECTION OPERATION

- Relation r

A	B	C
α	10	1
α	20	1
β	30	1
β	40	2

$\pi_{A,C}(r)$

A	C
α	1
α	1
β	1
β	2

=

A	C
α	1
β	1
β	2

UNION OPERATION

- Relations r, s

A	B
-----	-----

α	1
α	2
β	1

r

A	B
-----	-----

α	2
β	3

s

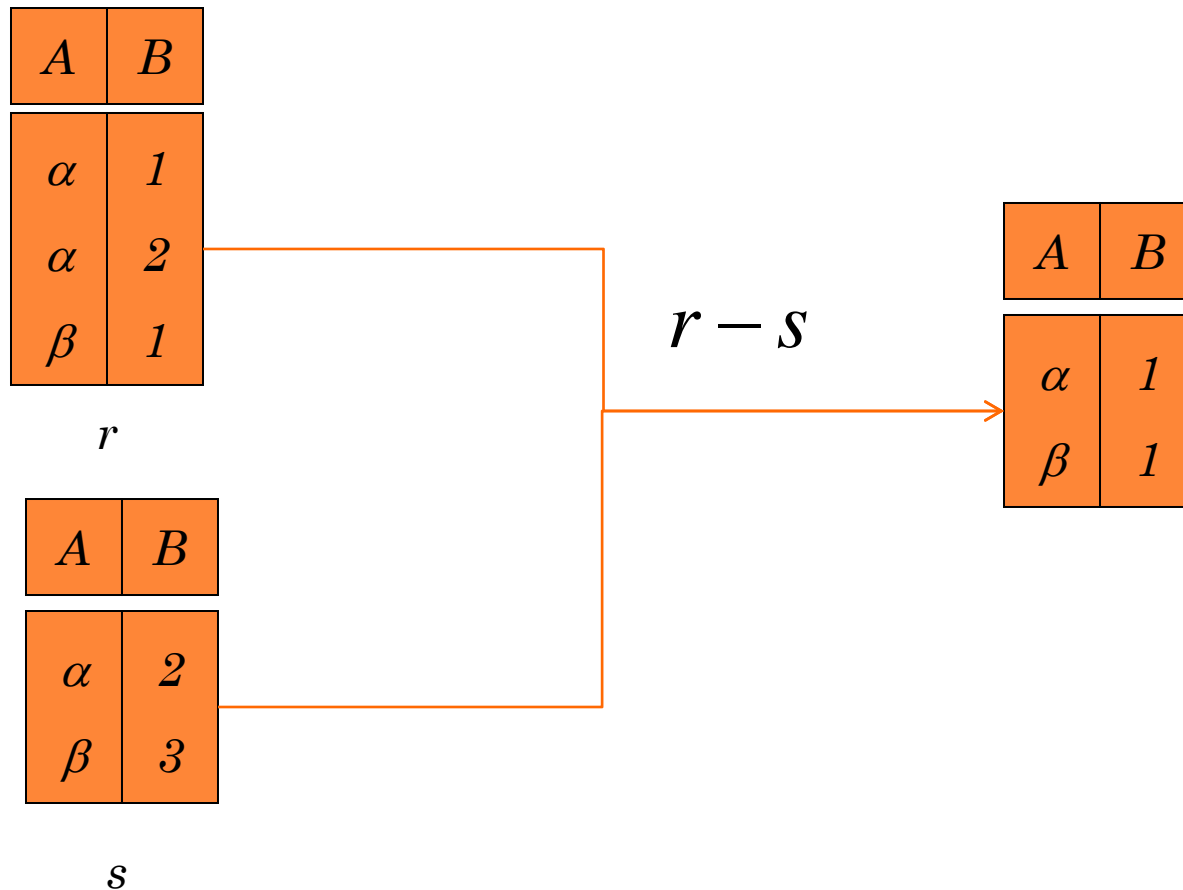
$r \cup s$

A	B
-----	-----

α	1
α	2
β	1
β	3

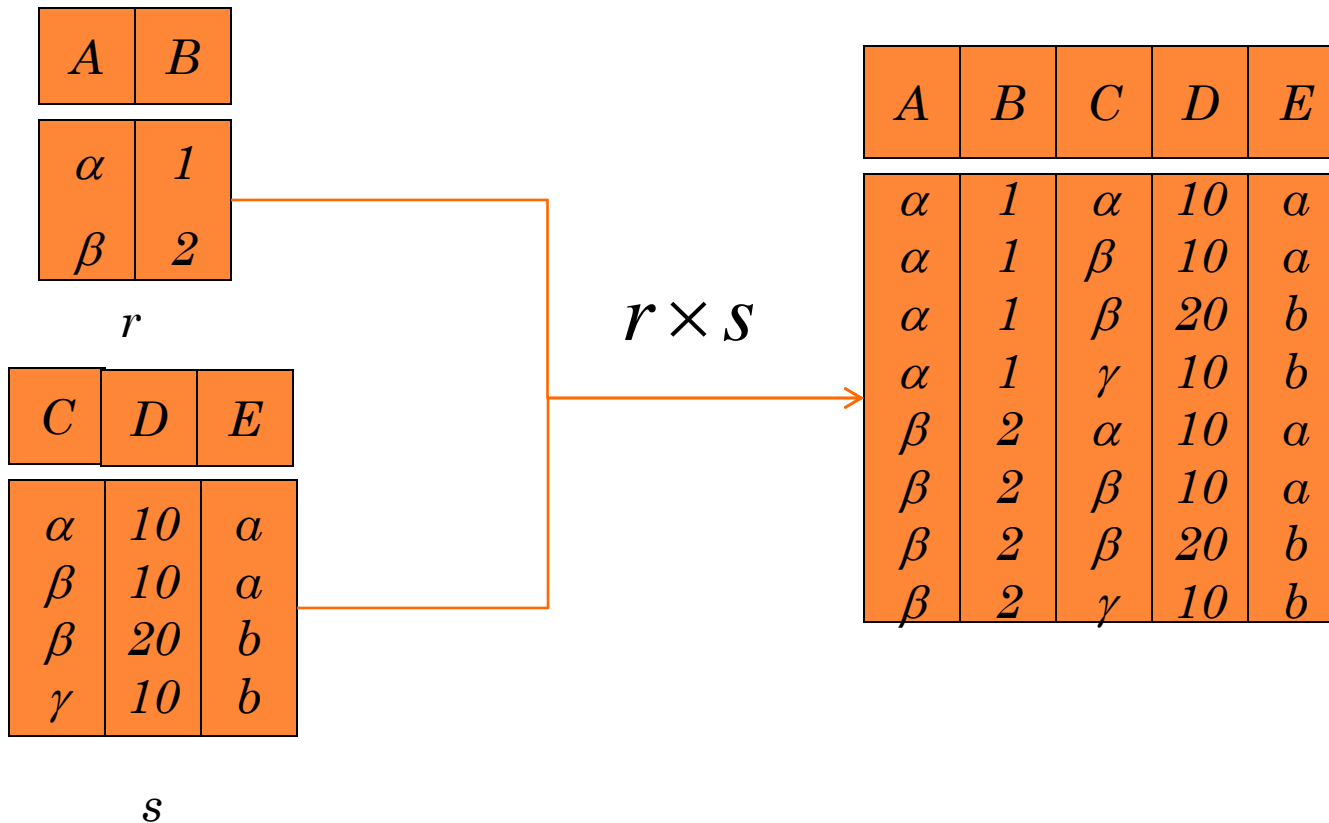
SET DIFFERENCE OPERATION

- Relations r, s



CARTESIAN PRODUCT OPERATION

- Relations r, s



RENAME OPERATION

- Allows us to name, and therefore to refer to, the results of relational-algebra expressions
- Allows us to refer to a relation by more than one name
- Example:

$$\rho_x(E)$$

returns the expression E under the name x

- If a relational-algebra expression E has arity n , then

$$\rho_{x(A_1, A_2, \dots, A_n)}(E)$$

returns the result of expression E under the name x , and with the attributes renamed to A_1, A_2, \dots, A_n .

BANKING EXAMPLE

branch (branch_name, branch_city, assets)

*customer (customer_name, customer_street,
customer_city)*


account (account_number, branch_name, balance)

loan (loan_number, branch_name, amount)

depositor (customer_name, account_number)

borrower (customer_name, loan_number)

BANK EXAMPLE



EXAMPLE QUERIES

- Find all loans of over Rs 2000

$$\sigma_{amount > 2000}(loan)$$

- Find the loan number for each loan of an amount greater than Rs 2000

$$\pi_{loan_no}(\sigma_{amount > 2000}(loan))$$

- Find the names of all customers who have a loan, an account, or both, from the bank

$$\pi_{customer_name}(depositor) \cup \pi_{customer_name}(borrower)$$

EXAMPLE QUERIES

- Find the names of all customers who have a loan at the Patliputra branch.

$$\pi_{customer_name}(\sigma_{(borrower.loan_no=loan.loan_no) \wedge (loan.branch=Patliputra)}(borrower \times loan))$$

- Find the names of all customers who have a loan at Patliputra branch but do not have an account at any branch of the bank.

$$\pi_{customer_name}(\sigma_{(borrower.loan_no=loan.loan_no) \wedge (loan.branch=Patliputra)}(borrower \times loan))$$


$$- \pi_{customer_name}(depositor)$$

SOME OTHER OPERATIONS

- Additional Operations
 - Set intersection
 - Natural join
 - Outer Join
 - Division
- The above operations can be expressed using basic operations we have seen earlier

SET INTERSECTION OPERATION

- Relations r, s



SET INTERSECTION OPERATION (CONTD.)

- Set intersection operation can be built using other basic operations
- How?


$$r \cap s = r - (r - s)$$

NATURAL JOIN OPERATION

- Cartesian product often requires a selection operation
- The selection operation most often requires that all attributes that are common to the relations are involved in the Cartesian product be equated
- Steps for natural join-
 1. Perform the Cartesian product of its two arguments
 2. Perform a selection forcing equality on those attributes that appear in both relational schemas
 3. Finally remove the duplicate attributes

NATURAL JOIN OPERATION

- Relations r, s



NATURAL JOIN OPERATION (CONTD.)

- A natural join operation can be rewritten as

$$r \bowtie s = \pi_{R \cup S} (\sigma_{r.A_1=s.A_1 \wedge r.A_2=s.A_2 \dots r.A_n=s.A_n} (r \times s))$$

- **Theta join** is a variant of natural join
- It is defined as

$$r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$$

ASSIGNMENT OPERATION

- It is convenient at times to write relational algebra expression by **assigning parts of it to temporary relation variables**
- The assignment operation works like assignment in programming languages
- We can rewrite $r \bowtie s$ as

temp1 \leftarrow $r \bowtie s$

temp2 $\leftarrow \sigma_{r.A_1=s.A_1 \wedge r.A_2=s.A_2 \dots r.A_n=s.A_n}(\text{temp1})$

result $\leftarrow \pi_{R \cup S}(\text{temp2})$

OUTER JOIN OPERATION

- An extension of the join operation that **avoids loss of information**
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join
- Uses *null* values:
 - *null* signifies that the value is unknown or does not exist
 - All comparisons involving *null* are **false** by definition.

DIFFERENT FORMS OF OUTER JOIN

- Left outer join \bowtie
 - Includes the tuples from the left relation that did not match with any tuples in the right relation
 - Pads the tuples with null values for all other attributes from the right relation
 - Adds them to the result of the natural join
- We can rewrite $r \bowtie s$ as

$$r \bowtie s \cup (r - \pi_R(r \bowtie s)) \times \{(\text{null}, \dots, \text{null})\}$$

Here the constant relation $\{(\text{null}, \dots, \text{null})\}$ is on schema S-R

DIFFERENT FORMS OF OUTER JOIN

- Similarly we can define-
- Right outer join $\bowtie\!\!\!\!\!\lrcorner$
- Full outer join \boxtimes

NULL VALUE

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist
- Aggregate functions simply ignore null values (as in SQL)
- For duplicate elimination and grouping, null is treated like any other value, and two nulls are assumed to be the same (as in SQL)

DIVISION OPERATION

- Notation: $r \div s$
- Suited to queries that include the phrase “*for all*”.
- Let r and s be relations on schemas R and S respectively where
 - $R = (A_1, \dots, A_m, B_1, \dots, B_n)$
 - $S = (B_1, \dots, B_n)$

The result of $r \div s$ is a relation on schema

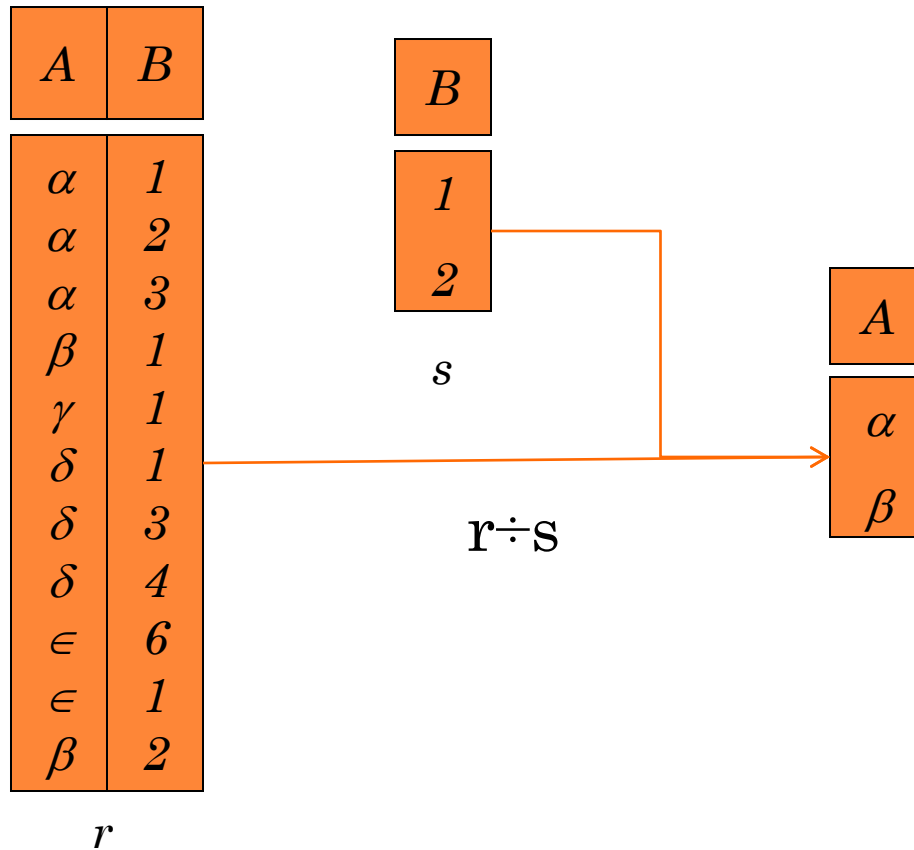
$$R - S = (A_1, \dots, A_m)$$

$$r \div s = \{ t \mid t \in \Pi_{R-S}(r) \wedge \forall u \in s (tu \in r) \}$$

Where tu means the concatenation of tuples t and u to produce a single tuple

DIVISION OPERATION (CONTD.)

- Relations r, s



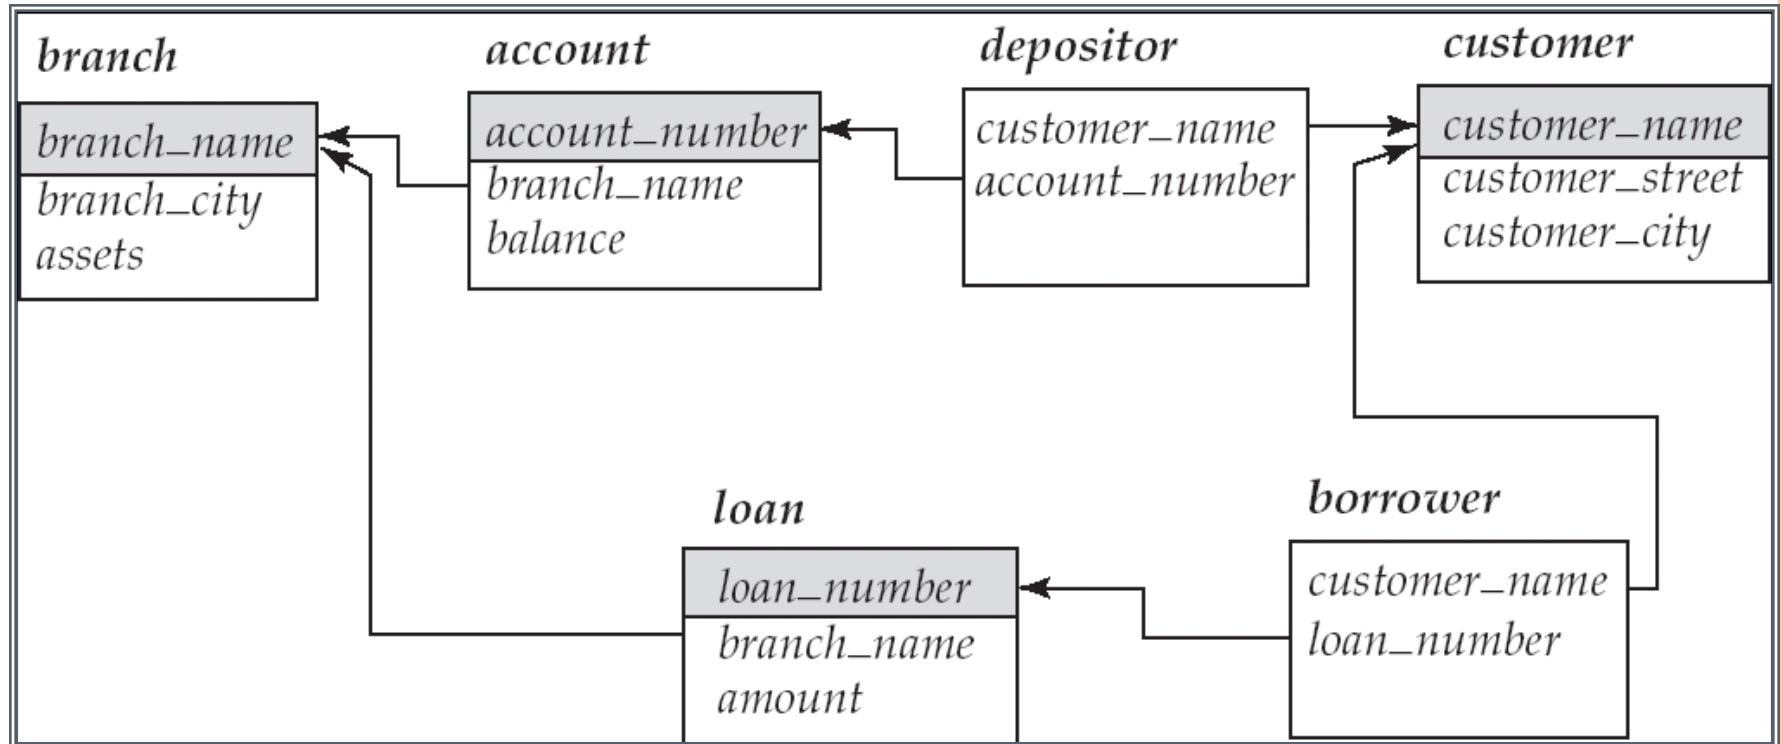
DIVISION OPERATION (CONTD.)

- Definition in terms of the basic algebra operation

Let $r(R)$ and $s(S)$ be relations, and let $S \subseteq R$

$$r \div s = \Pi_{R-S}(r) - \Pi_{R-S}((\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r))$$

BANK EXAMPLE



EXAMPLE QUERIES

- Find the names of all customers who have a loan and an account at bank.
 - $\Pi_{customer_name} (borrower) \cap \Pi_{customer_name} (depositor)$
- Find the name of all customers who have a loan at the bank and the loan amount
 - $\Pi_{customer_name, loan_number, amount} (borrower \bowtie loan)$

EXAMPLE QUERIES (CONTD.)

- Find the largest account balance in the bank
 - $\Pi_{balance}(account) - \Pi_{account.balance}(\sigma_{account.balance < d.balance}(account \times \rho_d account))$

EXAMPLE QUERIES (CONTD.)

- Find the name of customers who have an account at all the branches located in “Patna” city.

$$\Pi_{customer_name, branch_name} (depositor \bowtie account) \\ \div \Pi_{branch_name} (\sigma_{branch_city = \text{“Patna”}} (branch))$$

FEW MORE JOIN OPERATIONS

○ Semi join

- The left semi-join is similar to the natural join
- The result of this semi-join is the set of all tuples in r for which there is a tuple in s that is equal on their common attribute names
- $r \text{ semiJoin } s = \Pi_{A1, \dots, An}(r \text{ naturalJoin } s)$ where $R = \{A1, \dots, An\}$

○ Anti join

- It is similar to the natural join,
- but the result of an anti-join is only those tuples in r for which there is *no* tuple in s that is equal on their common attribute names
- $r \text{ antiJoin } s = r - (r \text{ semiJoin } s)$

EXTENDED RELATIONAL ALGEBRA OPERATIONS

- Generalized Projection
- Aggregation

GENERALIZED PROJECTION

- An **extension of projection** which allows operations such as **arithmetic and string functions** to be used in the projection list
- $\Pi_{F1,F2,\dots,Fn}(E)$
 - here $F1, F2, \dots, Fn$ is an arithmetic expression involving constants and attributes in the schema of E
- **Example:** $\Pi_{ID,name,dept_name,salary/12}(emp)$
- **Example:** $\Pi_{ID,(limit-balance) \text{ as } credit_available}(credit_info)$

AGGREGATION

- Aggregate functions take a collection of values and return a single value in result
 - E.g. *sum*, *avg*, *count*, *max*, *min*, etc.
- **Multisets**: the collections on which aggregate function operates can have multiple occurrences of a value; the order in which the values appear is not relevant
 - *g*_{sum(salary)}(instructor)



Caligraphic G

MODIFICATION OF THE DATABASE

- The content of the database may be modified using the following operations:
 - Deletion
 - Insertion
 - Updating
- All these operations are expressed using the assignment operator.


DELETION

- A delete request is expressed similarly to a query, except instead of displaying tuples to the user, the selected tuples are removed from the database
- Can delete only whole tuples; cannot delete values on only particular attributes
- A deletion is expressed in relational algebra by:

$$r \leftarrow r - E$$

where r is a relation and E is a relational algebra query

BANK EXAMPLE



DELETE EXAMPLE

○ Delete all account records from “Patliputra” branch

- $r1 \leftarrow \sigma_{branch_name="Patliputra"} (account \bowtie depositor)$
- $account \leftarrow account - \Pi_{account_number, branch_name, balance} (r1)$
- $depositor \leftarrow depositor - \Pi_{customer_name, account_number} (r1)$

○ Delete all loan records with amount in the range of 0 to 50

- $r2 \leftarrow \sigma_{amount \geq 0 \text{ and } amount \leq 50} (loan \bowtie borrower)$
- $loan \leftarrow loan - \Pi_{loan_no, branch_name, amount} (r2)$
- $borrower \leftarrow borrower - \Pi_{customer_name, loan_no} (r2)$

- Delete all accounts at branches located in city 'Gaya'

$r_1 \leftarrow \sigma_{branch_city = \text{"Gaya"}} (account \bowtie branch)$

$r_2 \leftarrow \Pi_{account_number, branch_name, balance} (r_1)$

$r_3 \leftarrow \Pi_{customer_name, account_number} (r_2 \bowtie depositor)$

$account \leftarrow account - r_2$

$depositor \leftarrow depositor - r_3$

INSERTION

- To insert data into a relation, we either:
 - specify a tuple to be inserted or
 - write a query whose result is a set of tuples to be inserted

- In relational algebra, an insertion is expressed by:

$$r \leftarrow r \cup E$$

where r is a relation and E is a relational algebra expression

- The insertion of a single tuple is expressed by letting E be a constant relation containing one tuple

INSERT EXAMPLE

- Insert information in the database specifying that Sumit has Rs 1200 in account A-973 at the Patliputra branch

$account \leftarrow account \cup \{("A-973", "Patliputra", 1200)\}$

$depositor \leftarrow depositor \cup \{("Sumit", "A-973")\}$

- Provide as a gift for all loan customers in the Patliputra branch, a Rs 200 savings account. Let the loan number serve as the account number for the new savings account

$r_1 \leftarrow (\sigma_{branch_name = "Patliputra"} (borrower \bowtie loan))$

$r_2 \leftarrow (\Pi_{loan_number, branch_name} (r_1))$

$account \leftarrow account \cup (r_2 \times \{(200)\})$

$depositor \leftarrow depositor \cup \Pi_{customer_name, loan_number} (r_1)$

UPDATING

- A mechanism to **change a value in a tuple without changing *all* values in the tuple**
- Use the generalized projection operator to do this task

$$r \leftarrow \Pi_{F_1, F_2, \dots, F_i}(r)$$

Each F_i is either

- the i^{th} attribute of r , if the i^{th} attribute is not updated, or,
- if the attribute is to be updated F_i is an expression, involving only constants and the attributes of r , which gives the new value for the attribute

UPDATE EXAMPLE

- Make interest payments by increasing all balances by 5 percent
- $account \leftarrow \Pi_{account_number, branch_name, balance * 1.05}(account)$
- Pay all accounts with balances over Rs1,00,000 a six percent interest and pay all others five percent
- $account \leftarrow \Pi_{account_number, branch_name, balance * 1.06}(\sigma_{balance > 100000}(account)) \cup \Pi_{account_number, branch_name, balance * 1.05}(\sigma_{balance \leq 100000}(account))$

TUPLE RELATIONAL CALCULUS

TUPLE RELATIONAL CALCULUS

- A **nonprocedural** query language, where each query is of the form

$$\{t \mid P(t)\}$$

- It is the set of all tuples t such that predicate P is true for t
- t is a *tuple variable*, $t[A]$ denotes the value of tuple t on attribute A
- $t \in r$ denotes that tuple t is in relation r
- P is a *formula* similar to that of the predicate calculus


PREDICATE CALCULUS FORMULA

- ❑ Set of attributes and constants
- ❑ Set of comparison operators: (e.g., $<$, \leq , $=$, \neq , $>$, \geq)
- ❑ Set of connectives: *and* (\wedge), *or* (\vee), *not* (\neg)
- ❑ Implication (\Rightarrow): $x \Rightarrow y$, **if x is true, then y is true**

$$x \Rightarrow y \equiv \neg x \vee y$$

- ❑ Set of quantifiers:
 - ▶ $\exists t \in r (Q(t)) \equiv$ “there exists” a tuple in t in relation r such that predicate $Q(t)$ is true
 - ▶ $\forall t \in r (Q(t)) \equiv Q(t)$ is true “for all” tuples t in relation r

BANK EXAMPLE



EXAMPLE QUERIES

- Find the *loan_number*, *branch_name*, and *amount* for loans of over Rs1200
 - $\{t \mid t \in \text{loan} \wedge t[\text{amount}] > 1200\}$
- Find the loan number for each loan of an amount greater than Rs1200
 - $\{t \mid \exists s \in \text{loan} (t[\text{loan_number}] = s[\text{loan_number}] \wedge s[\text{amount}] > 1200)\}$
- Find the names of all customers having a loan, an account, or both at the bank
 - $\{t \mid \exists s \in \text{borrower} (t[\text{customer_name}] = s[\text{customer_name}]) \vee \exists u \in \text{depositor} (t[\text{customer_name}] = u[\text{customer_name}])\}$

EXAMPLE QUERIES (2)

- Find the names of all customers who have a loan and an account at the bank

- $$\{t \mid \exists s \in \text{borrower} (t [\text{customer_name}] = s [\text{customer_name}]) \\ \wedge \exists u \in \text{depositor} (t [\text{customer_name}] = u [\text{customer_name}])\}$$

- Find the names of all customers having a loan at Patliputra branch

- $$\{t \mid \exists s \in \text{borrower} (t [\text{customer_name}] = s [\text{customer_name}] \\ \wedge \exists u \in \text{loan} (u [\text{branch_name}] = \text{"Patliputra"} \\ \wedge u [\text{loan_number}] = s [\text{loan_number}]))\}$$

EXAMPLE QUERIES (3)

- Find the names of all customers who have a loan at the Patliputra branch, but no account at any branch of the bank
 - $\{t \mid \exists s \in \text{borrower } (t [\text{customer_name}] = s [\text{customer_name}]$
 $\wedge \exists u \in \text{loan } (u [\text{branch_name}] = \text{"Patliputra"}$
 $\wedge u [\text{loan_number}] = s [\text{loan_number}]))$
 $\wedge \textbf{not } \exists v \in \text{depositor } (v [\text{customer_name}] =$
 $t [\text{customer_name}])\}$

EXAMPLE QUERIES (4)

- Find the names of all customers having a loan from Patliputra branch, and the cities in which they live
 - $\{t \mid \exists s \in \text{loan} (s [\text{branch_name}] = \text{"Patliputra"} \\ \wedge \exists u \in \text{borrower} (u [\text{loan_number}] = s [\text{loan_number}] \\ \wedge t [\text{customer_name}] = u [\text{customer_name}] \\ \wedge \exists v \in \text{customer} (u [\text{customer_name}] = v \\ [\text{customer_name}] \\ \wedge t [\text{customer_city}] = v \\ [\text{customer_city}])))))\}$

EXAMPLE QUERIES (5)

- Find the names of all customers who have an account at all branches located in `branch_city = "Dhanbad"`:
 - $\{t \mid \forall u \in \text{branch } (u[\text{branch_city}] = \text{"Dhanbad"} \Rightarrow \exists v \in \text{account } (v[\text{branch_name}] = u[\text{branch_name}] \wedge \exists w \in \text{depositor } (w[\text{account_number}] = v[\text{account_number}] \wedge (t[\text{customer_name}] = w[\text{customer_name}])))\}$

SAFETY OF EXPRESSION

- It is possible to write tuple calculus expressions that generate **infinite relation**.
 - For example, $\{ t \mid \neg t \in r \}$ results in an infinite relation
- To guard against the problem, we restrict the set of allowable expressions to **safe expressions**

SAFE EXPRESSION

- Let's consider an expression $\{t \mid P(t)\}$ in the tuple relational calculus
 - $\text{dom}(P)$: the set of all values referenced by P
 - They include the values mentioned in P itself, as well as values appear in a tuple of a relation mentioned in P or constants that appear in P
 - $\text{dom}(t \mid t \in r(t[A] = 5))$ is the set containing 5 as well as the values appearing in any attribute of any tuple in r
- The expression is said to be *safe* if every component of t appears from the $\text{dom}(P)$
 - E.g. $\{t \mid t \in r(t[A] = 5 \vee \mathbf{true})\}$ is not safe --- as the result is an infinite set and some of the values may not appear in any relation or tuples or constants in P .

DOMAIN RELATIONAL CALCULUS


DOMAIN RELATIONAL CALCULUS

- Another **nonprocedural** query language equivalent in power to the tuple relational calculus
- Forms the basis of widely used **QBE** (Query By Example) language
- Each query is an expression of the form:

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid P(x_1, x_2, \dots, x_n) \}$$

- x_1, x_2, \dots, x_n represent domain variables
- P represents a formula similar to that of the predicate calculus

BANK EXAMPLE



EXAMPLE QUERIES

- Find the *loan_number*, *branch_name*, and *amount* for loans of over Rs1200
 - $\{ \langle l, b, a \rangle \mid \langle l, b, a \rangle \in \text{loan} \wedge a > 1200 \}$
- Find the names of all customers who have a loan of over Rs1200
 - $\{ \langle c \rangle \mid \exists l, b, a (\langle c, l \rangle \in \text{borrower} \wedge \langle l, b, a \rangle \in \text{loan} \wedge a > 1200) \}$
- Find the names of all customers who have a loan from the Patliputra branch and the loan amount:
 - ▶ $\{ \langle c, a \rangle \mid \exists l (\langle c, l \rangle \in \text{borrower} \wedge \exists b (\langle l, b, a \rangle \in \text{loan} \wedge b = \text{"Patliputra"})) \}$
 - ▶ $\{ \langle c, a \rangle \mid \exists l (\langle c, l \rangle \in \text{borrower} \wedge \langle l, \text{"Patliputra"}, a \rangle \in \text{loan}) \}$

EXAMPLE QUERIES (2)

- Find the names of all customers having a loan, an account, or both at the Patliputra branch:
 - $\{ \langle c \rangle \mid \exists l (\langle c, l \rangle \in \text{borrower} \wedge \exists b, a (\langle l, b, a \rangle \in \text{loan} \wedge b = \text{"Patliputra"}) \vee \exists a (\langle c, a \rangle \in \text{depositor} \wedge \exists b, n (\langle a, b, n \rangle \in \text{account} \wedge b = \text{"Patliputra"}))) \}$

EXAMPLE QUERIES (3)

- Find the names of all customers who have an account at all branches located in Dhanbad:
- $\{ \langle c \rangle \mid \forall x, y, z (\langle x, y, z \rangle \in \text{branch} \wedge y = \text{"Dhanbad"}) \Rightarrow \exists a, b (\langle a, x, b \rangle \in \text{account} \wedge \langle c, a \rangle \in \text{depositor}) \}$

SAFETY OF EXPRESSION

The expression:

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid P(x_1, x_2, \dots, x_n) \}$$

is safe if -

all values that appear in tuples of the expression
are values from *dom* (*P*)

EXPRESSIVE POWER OF LANGUAGES

- The following languages are equivalent
 - The basic relational algebra (without extended relational algebra operation)
 - The tuple relational calculus (restricted to safe expression)
 - The domain relational calculus (restricted to safe expression)